

# Exploring 3D-aware Latent Spaces for Efficiently Learning Numerous Scenes

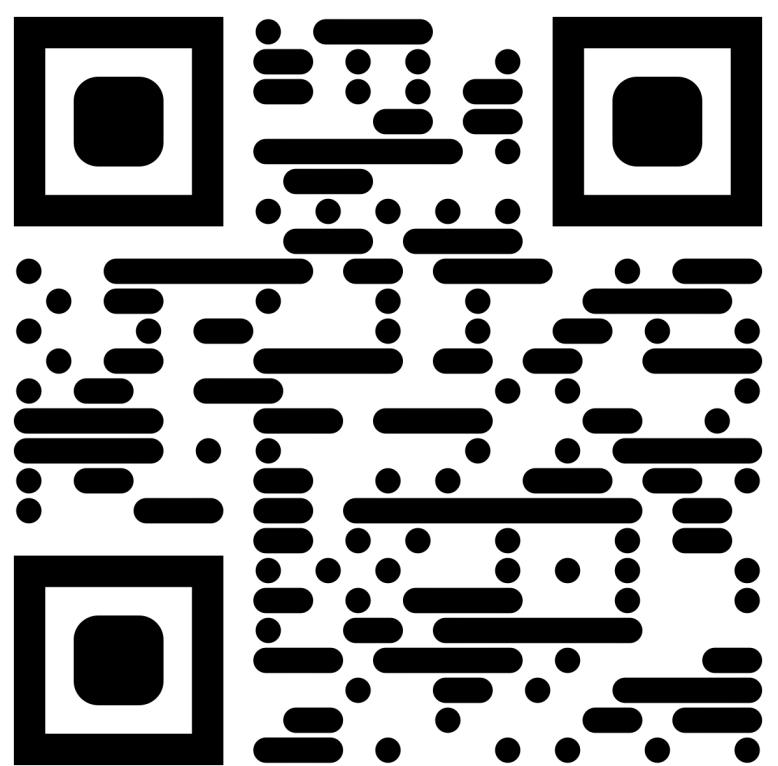
Antoine Schnepf<sup>\*1,3</sup>, Karim Kassab<sup>\*1,2</sup>,  
Jean-Yves Franceschi<sup>1</sup>, Laurent Caraffa<sup>2</sup>, Flavian Vasile<sup>1</sup>, Jeremie Mary<sup>1</sup>,  
Andrew Comport<sup>3</sup>, Valérie Gouet-Brunet<sup>2</sup>

<sup>\*</sup> Equal Contributions

<sup>1</sup> Criteo AI Lab, Paris, France

<sup>2</sup> LASTIG, Université Gustave Eiffel, IGN-ENSG, F-94160 Saint-Mandé

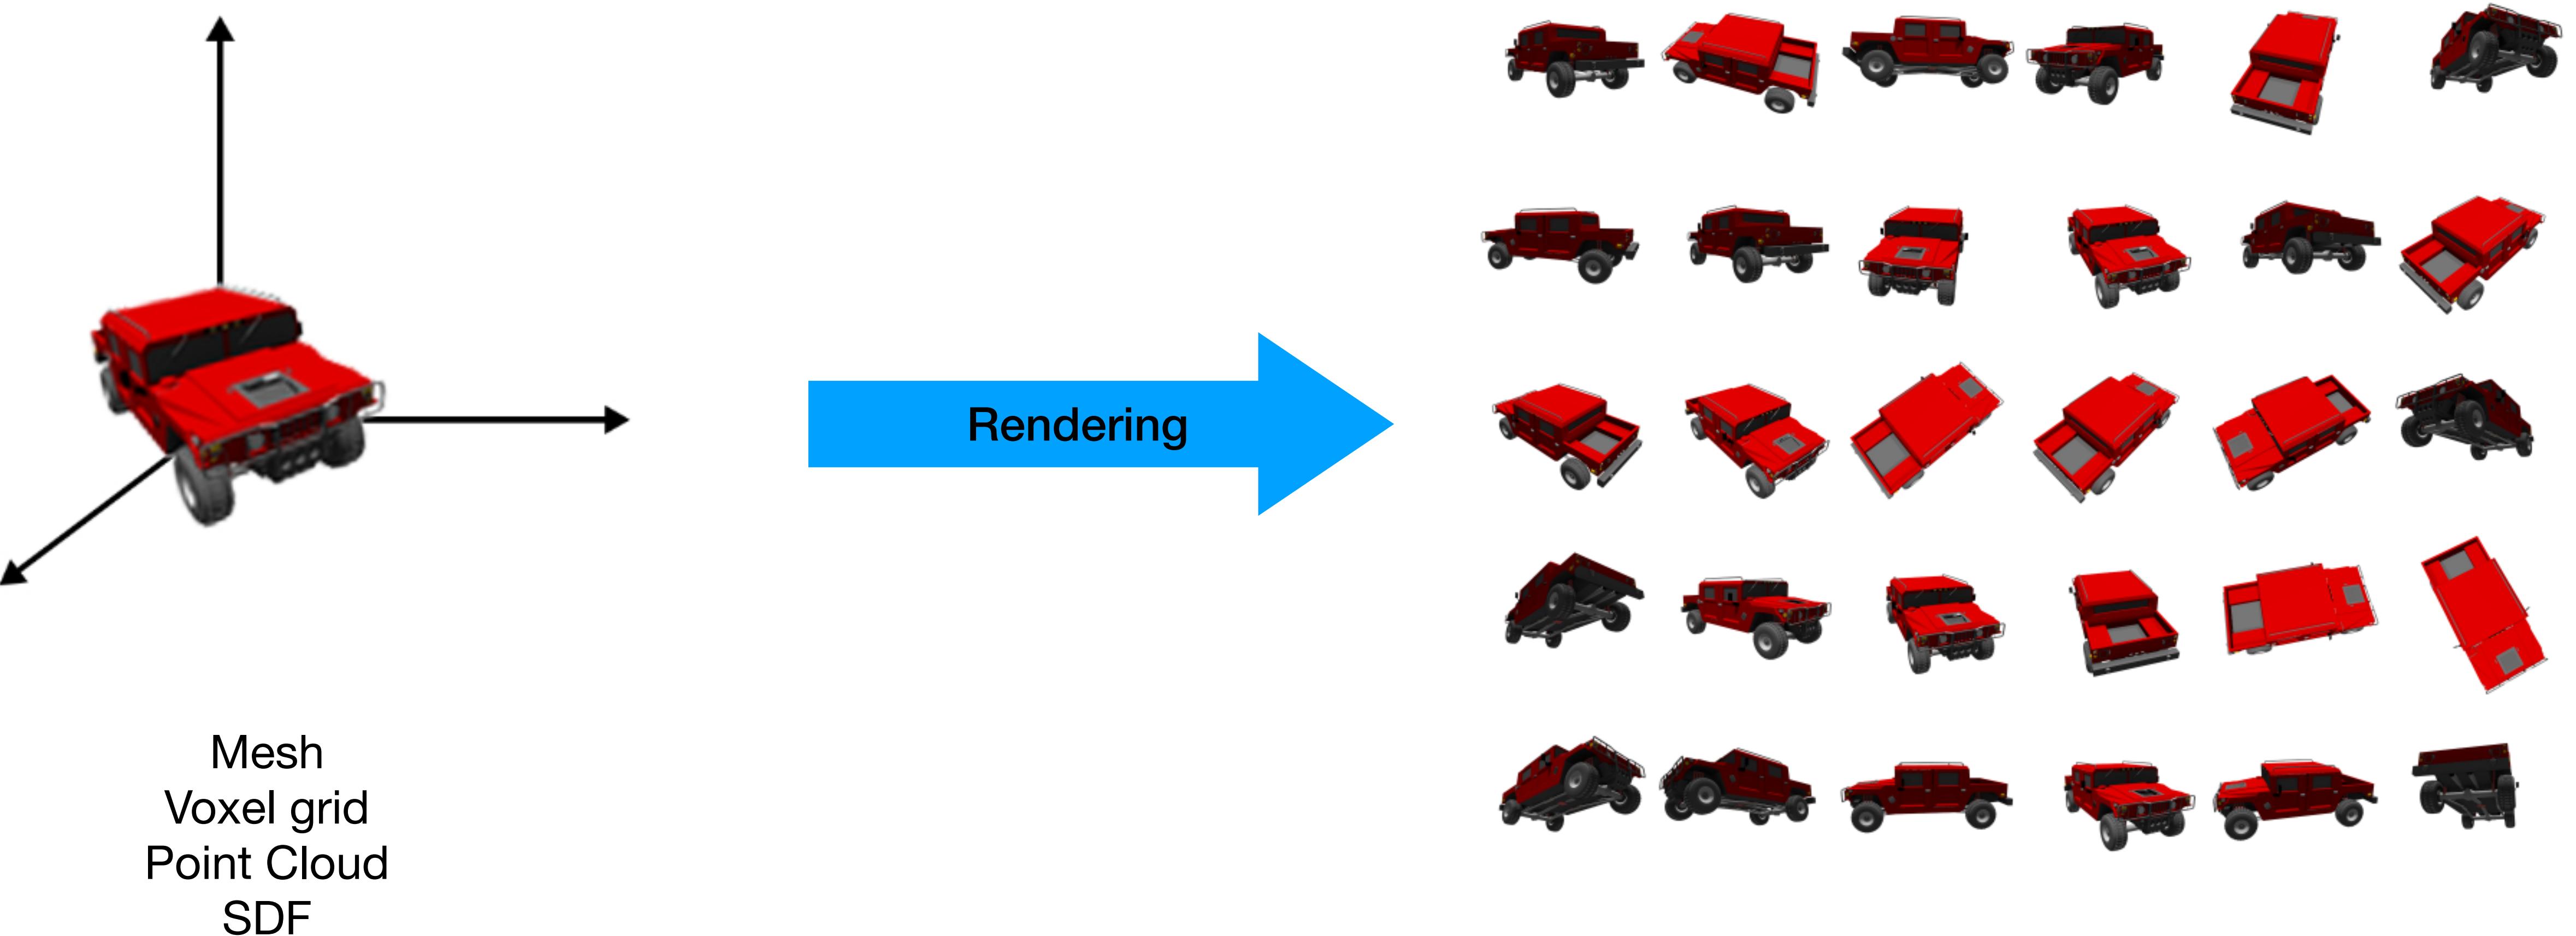
<sup>3</sup> Université Côte d'Azur, CNRS, I3S, France



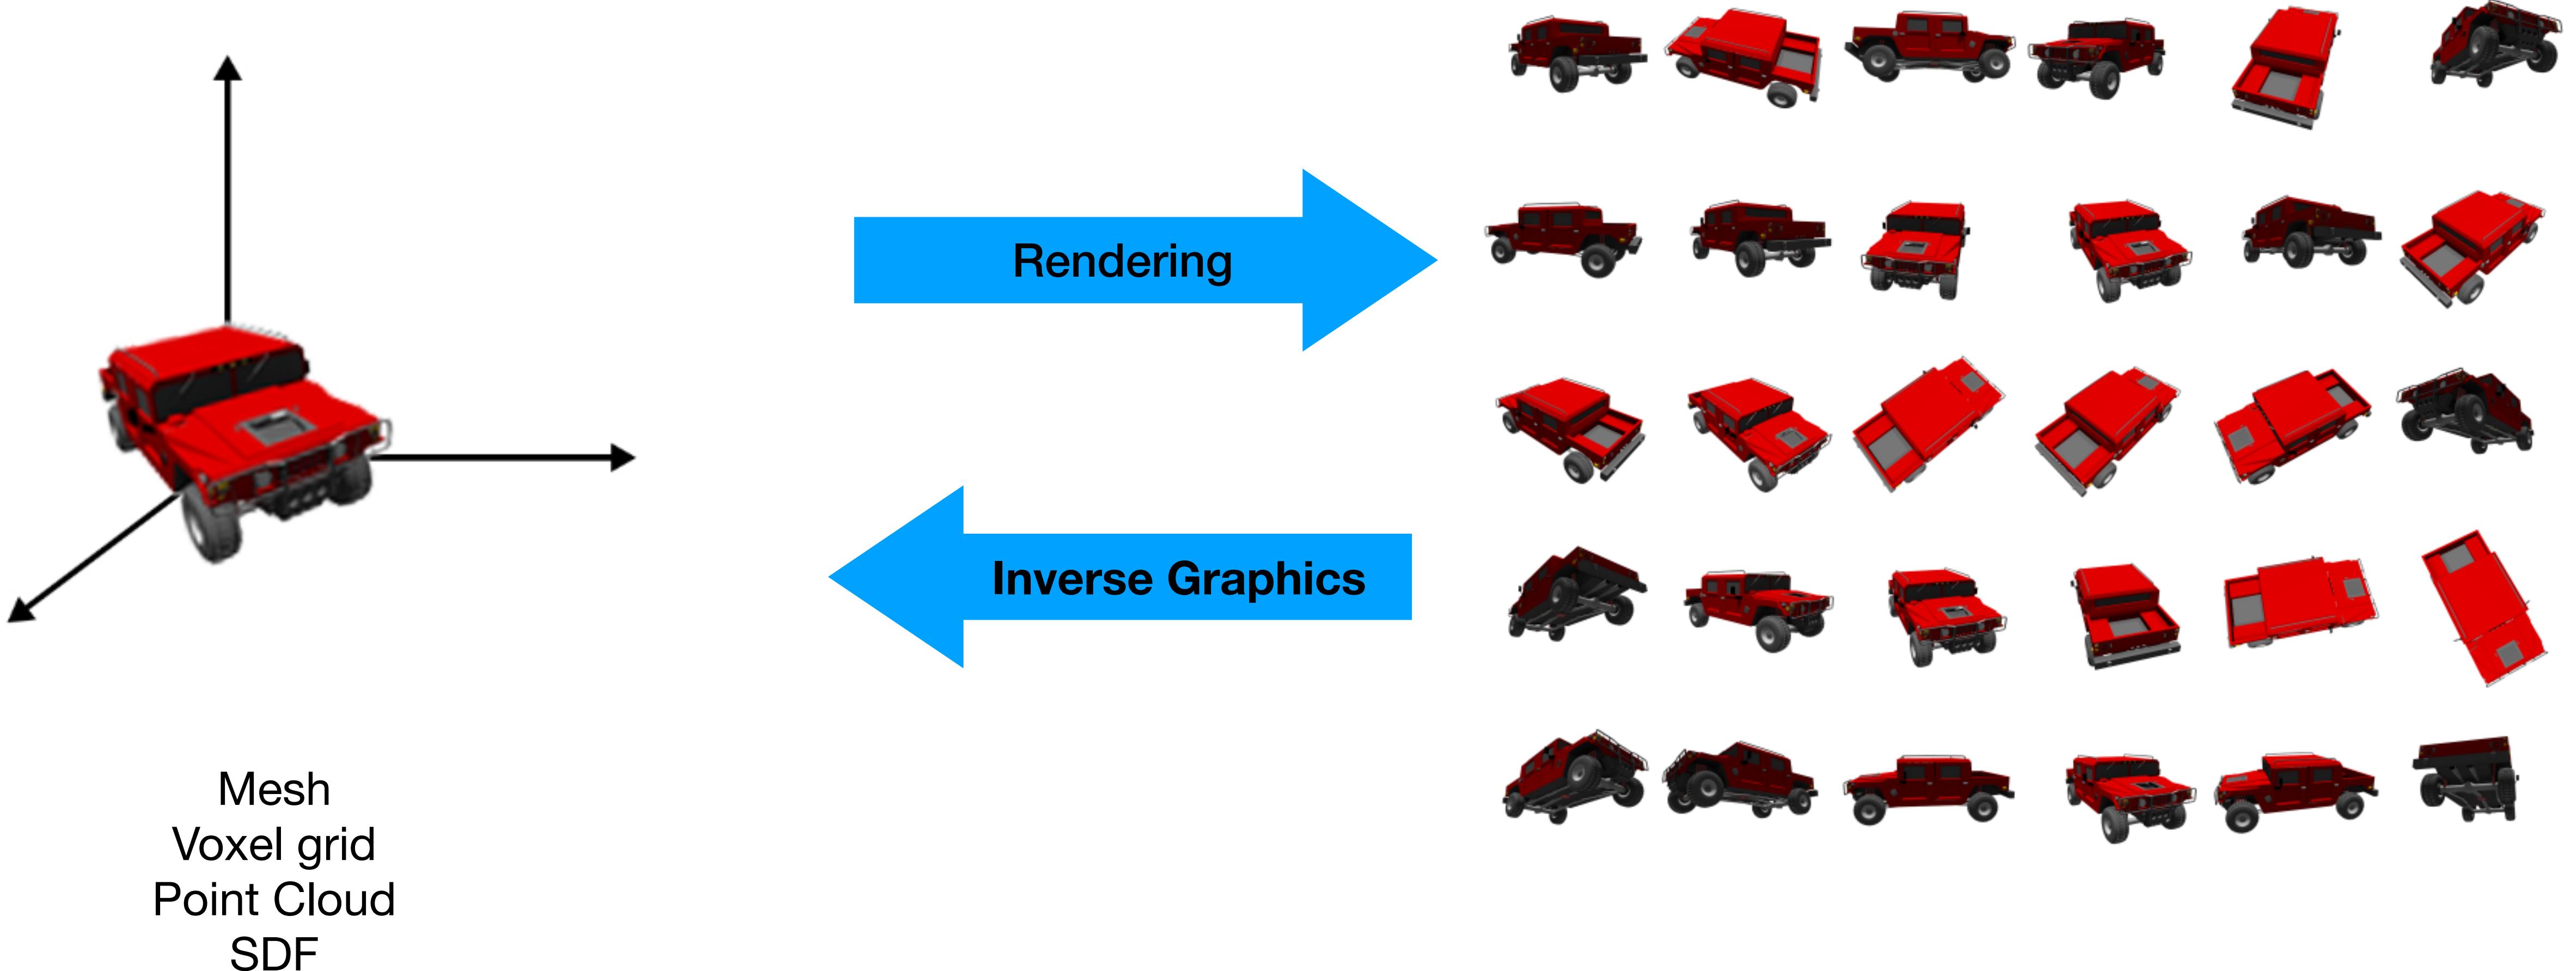
3da-ae.github.io

# Pre-requisites

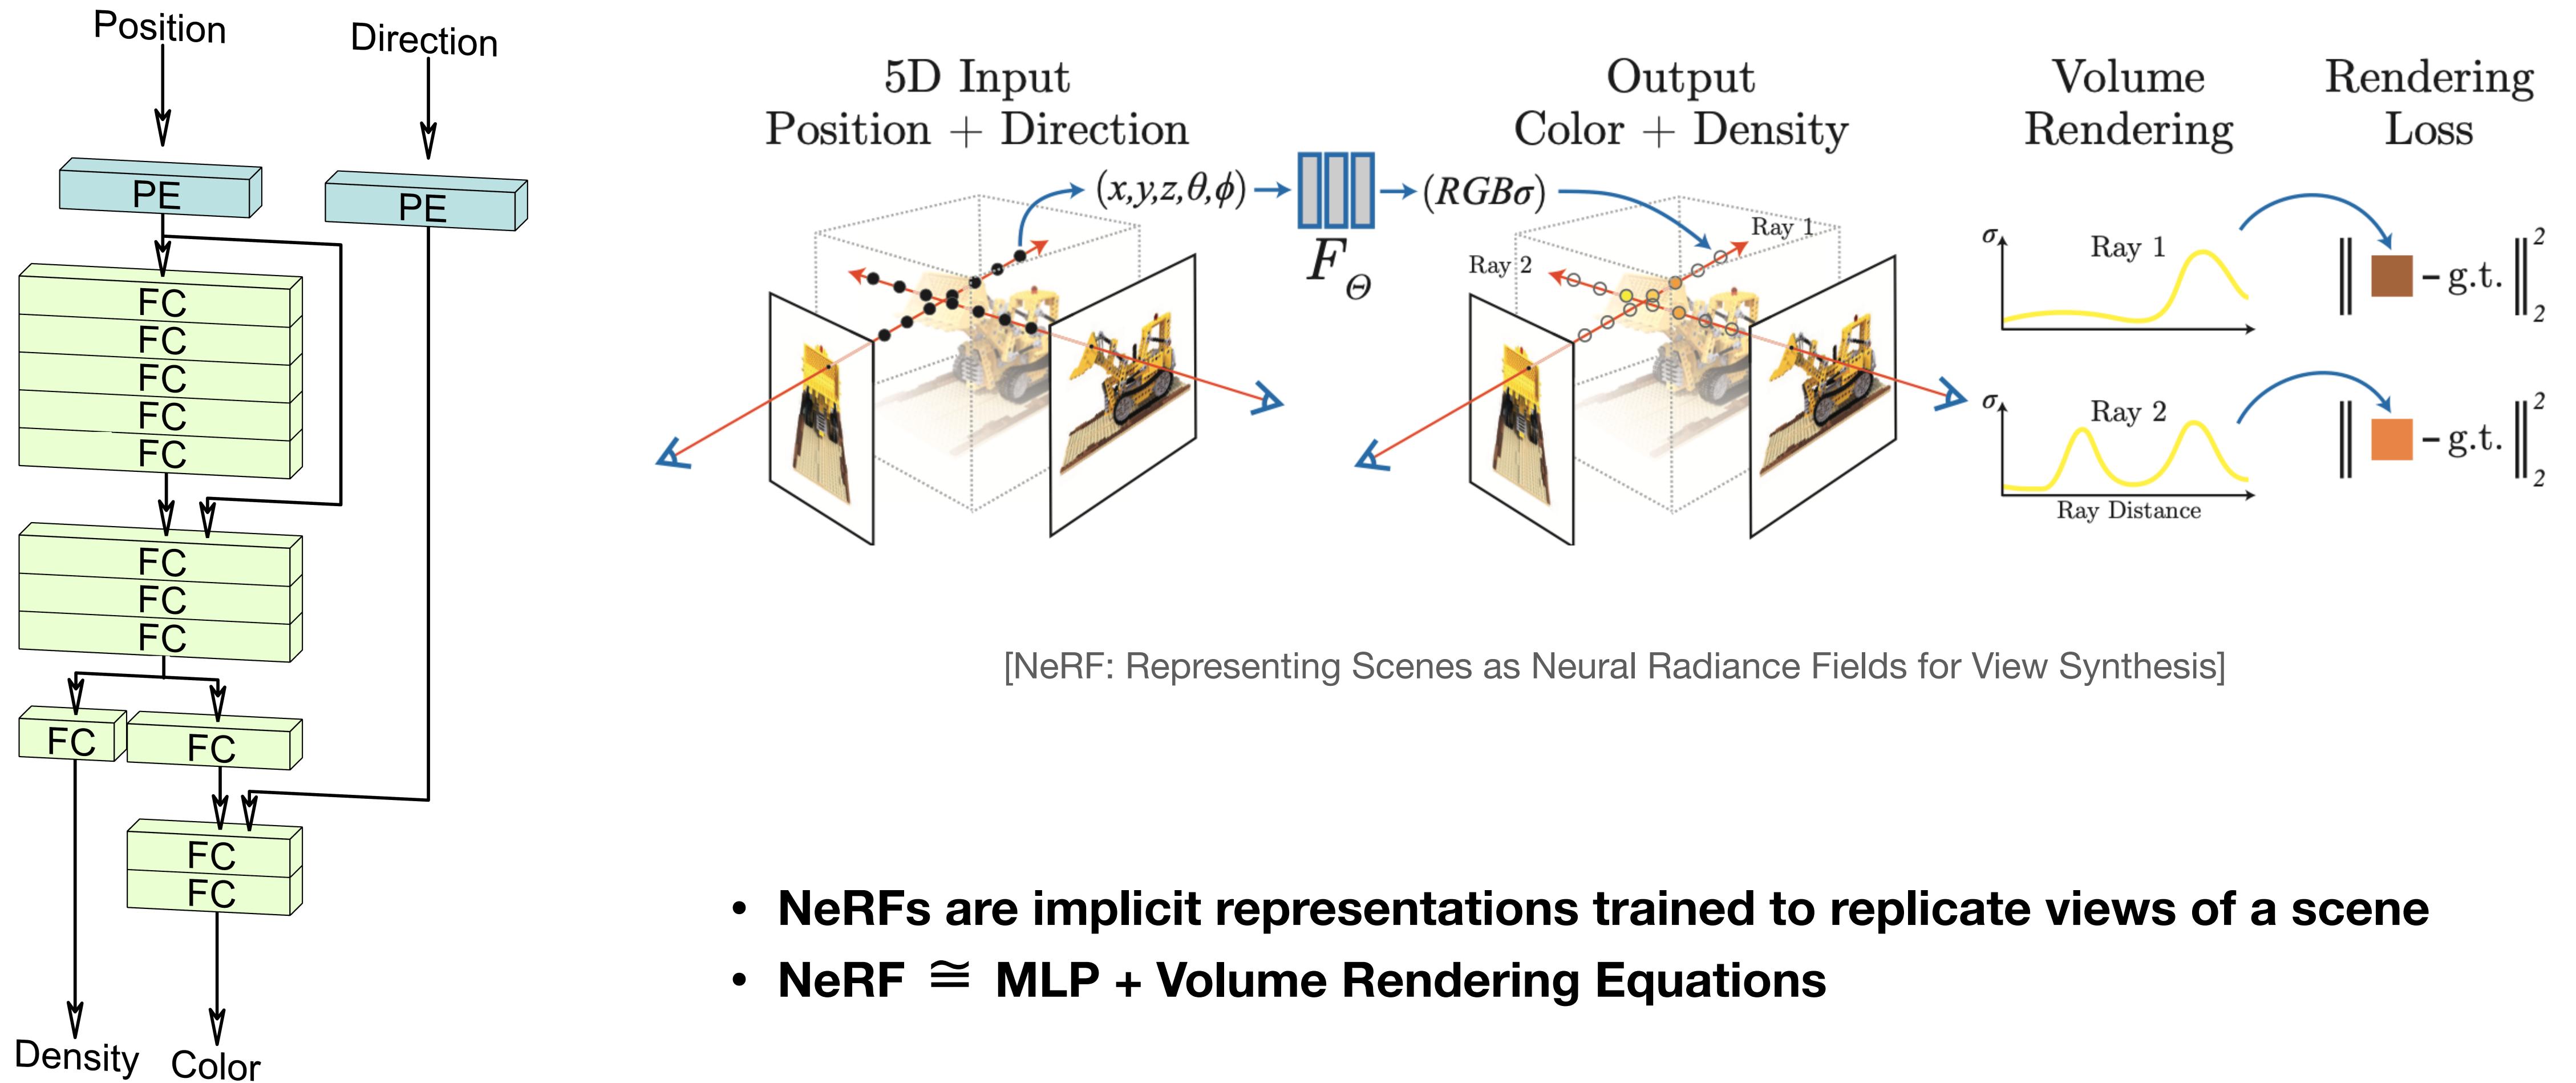
# The inverse graphics problem



# The inverse graphics problem

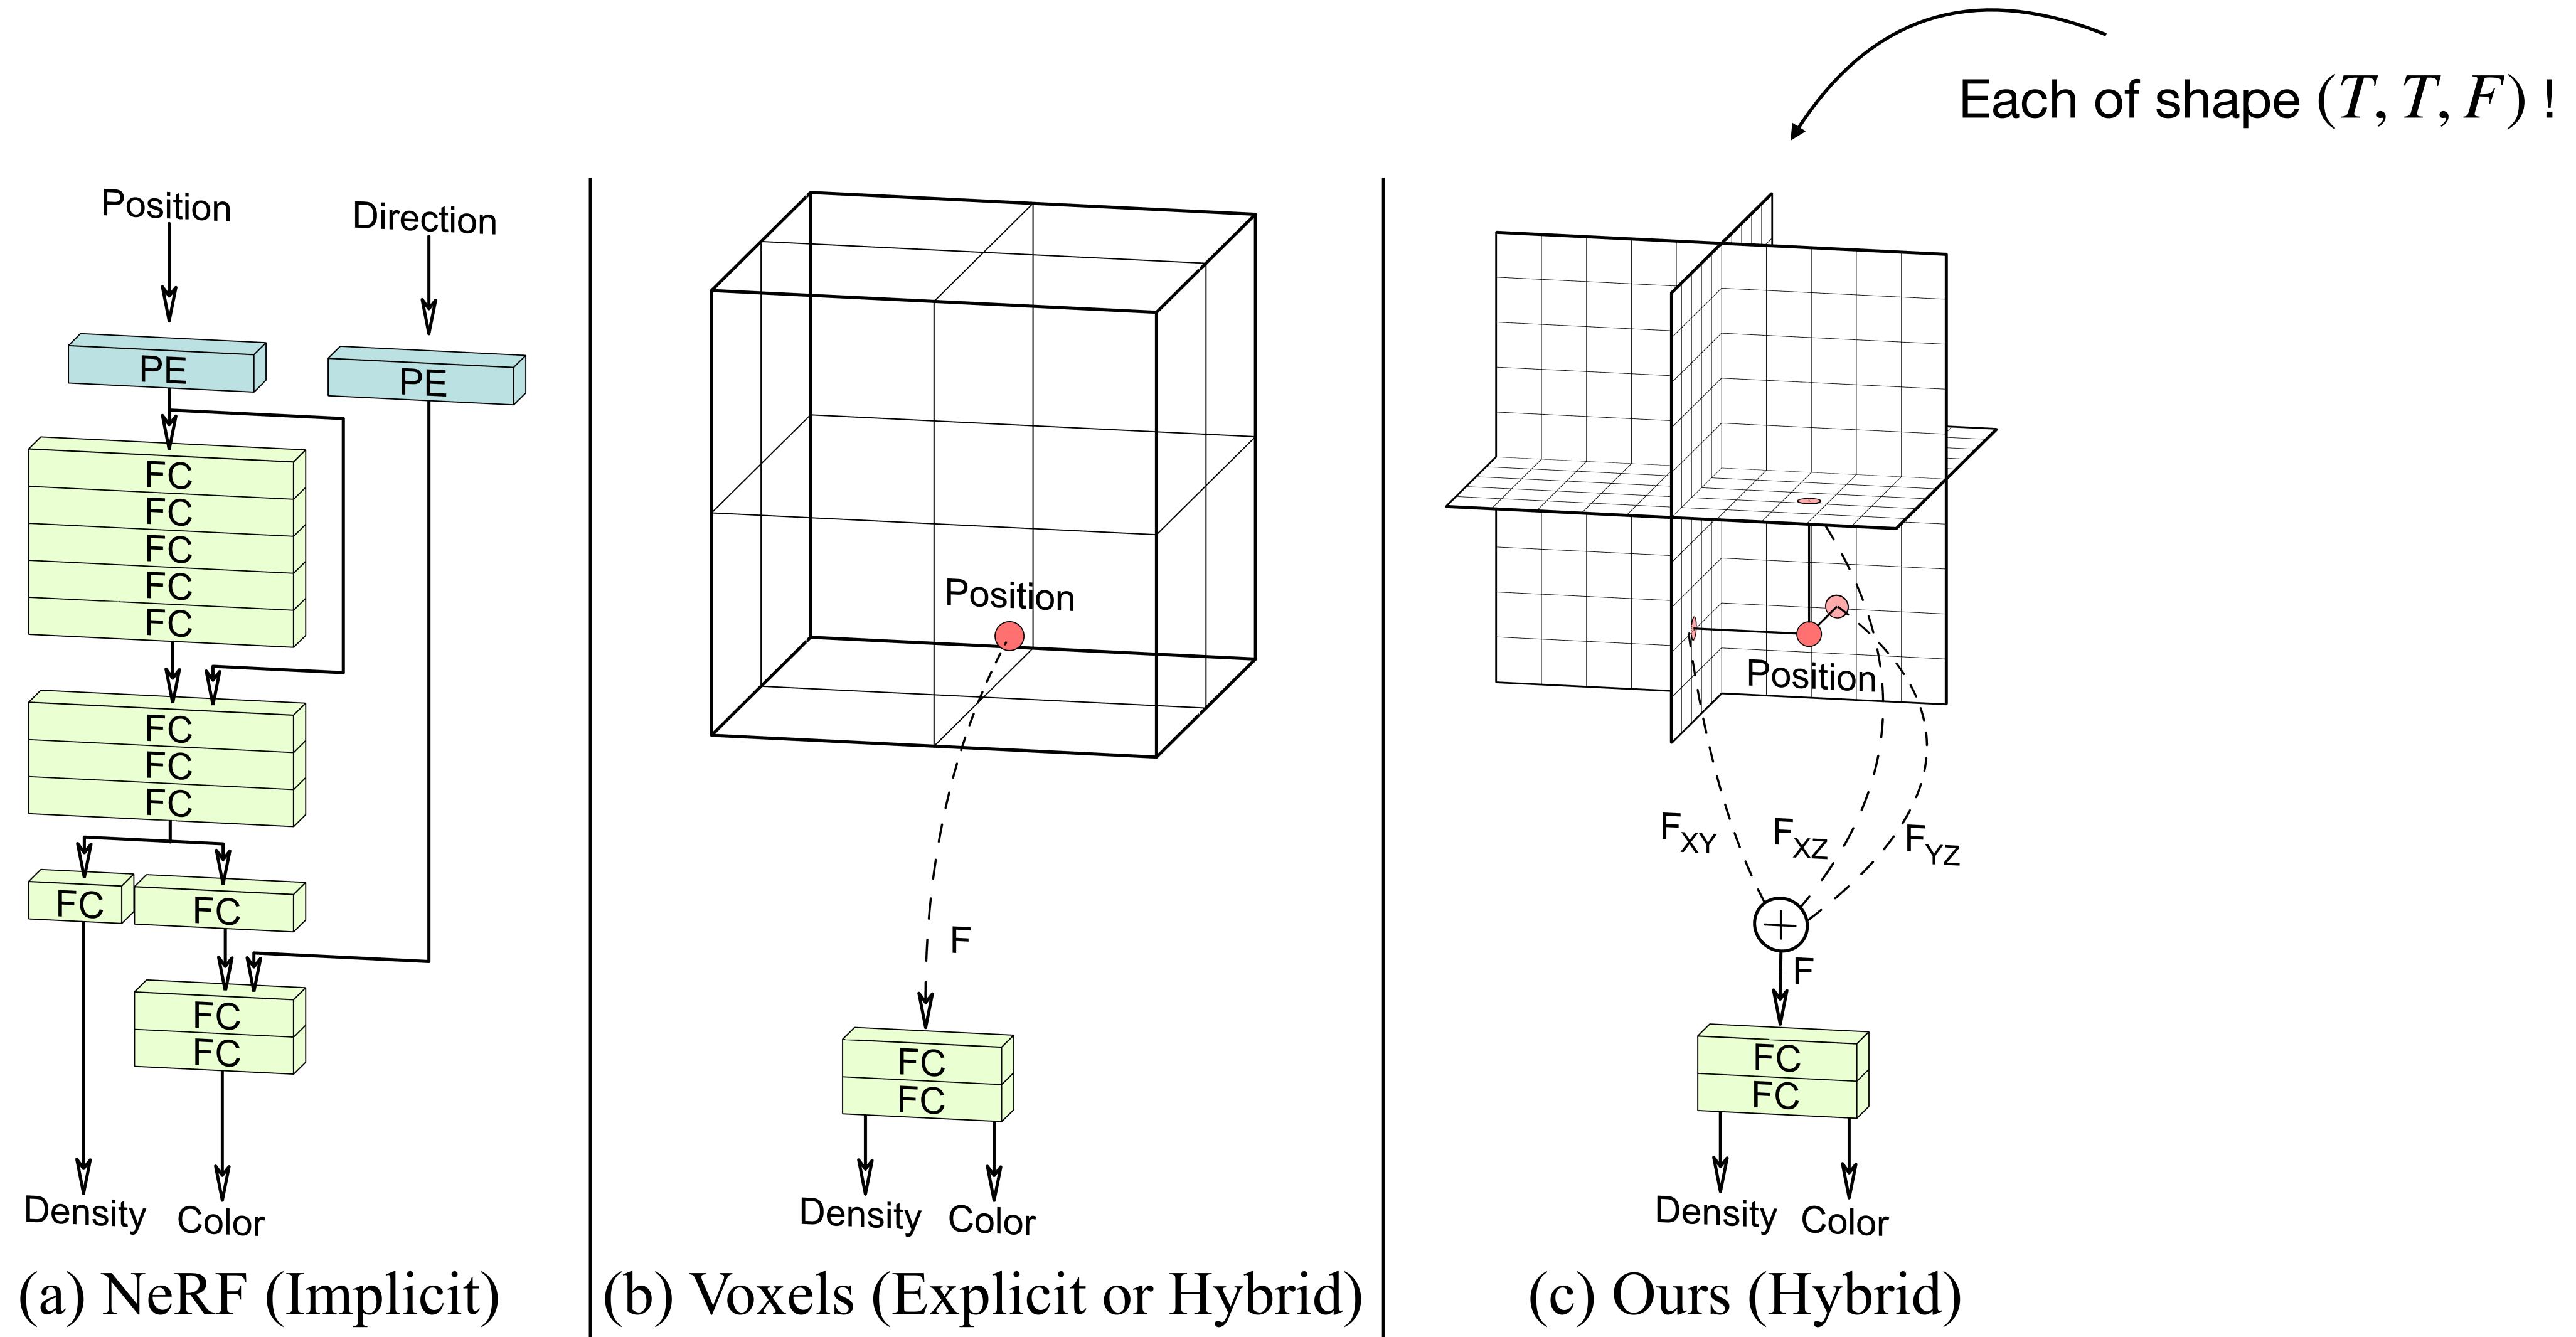


# Neural Radiance Fields



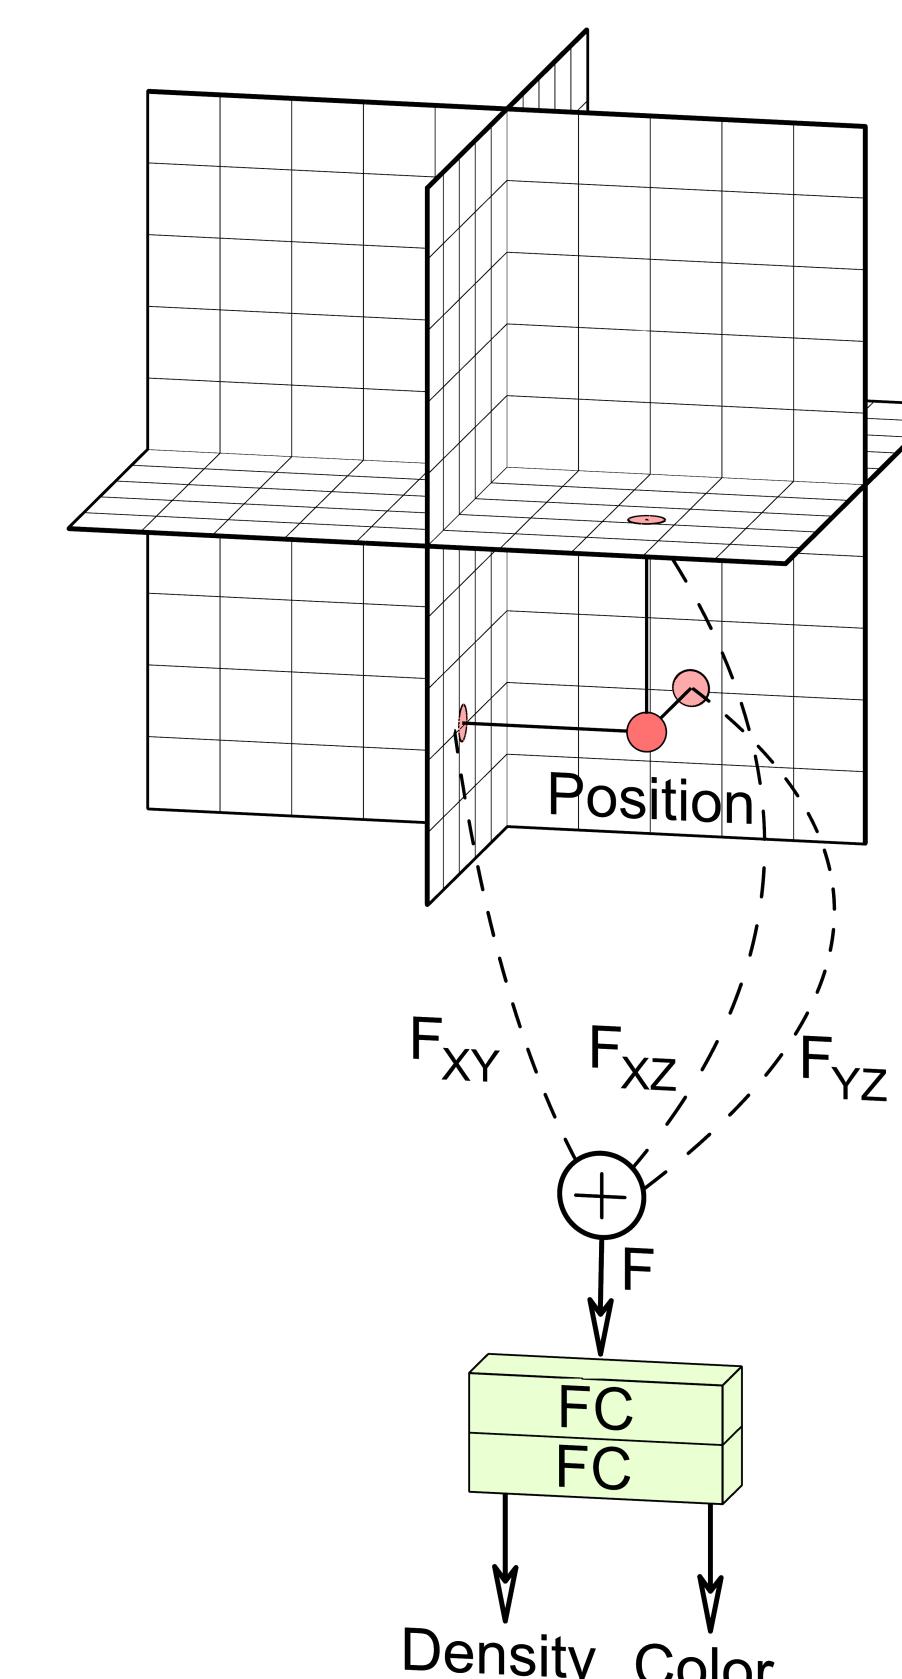
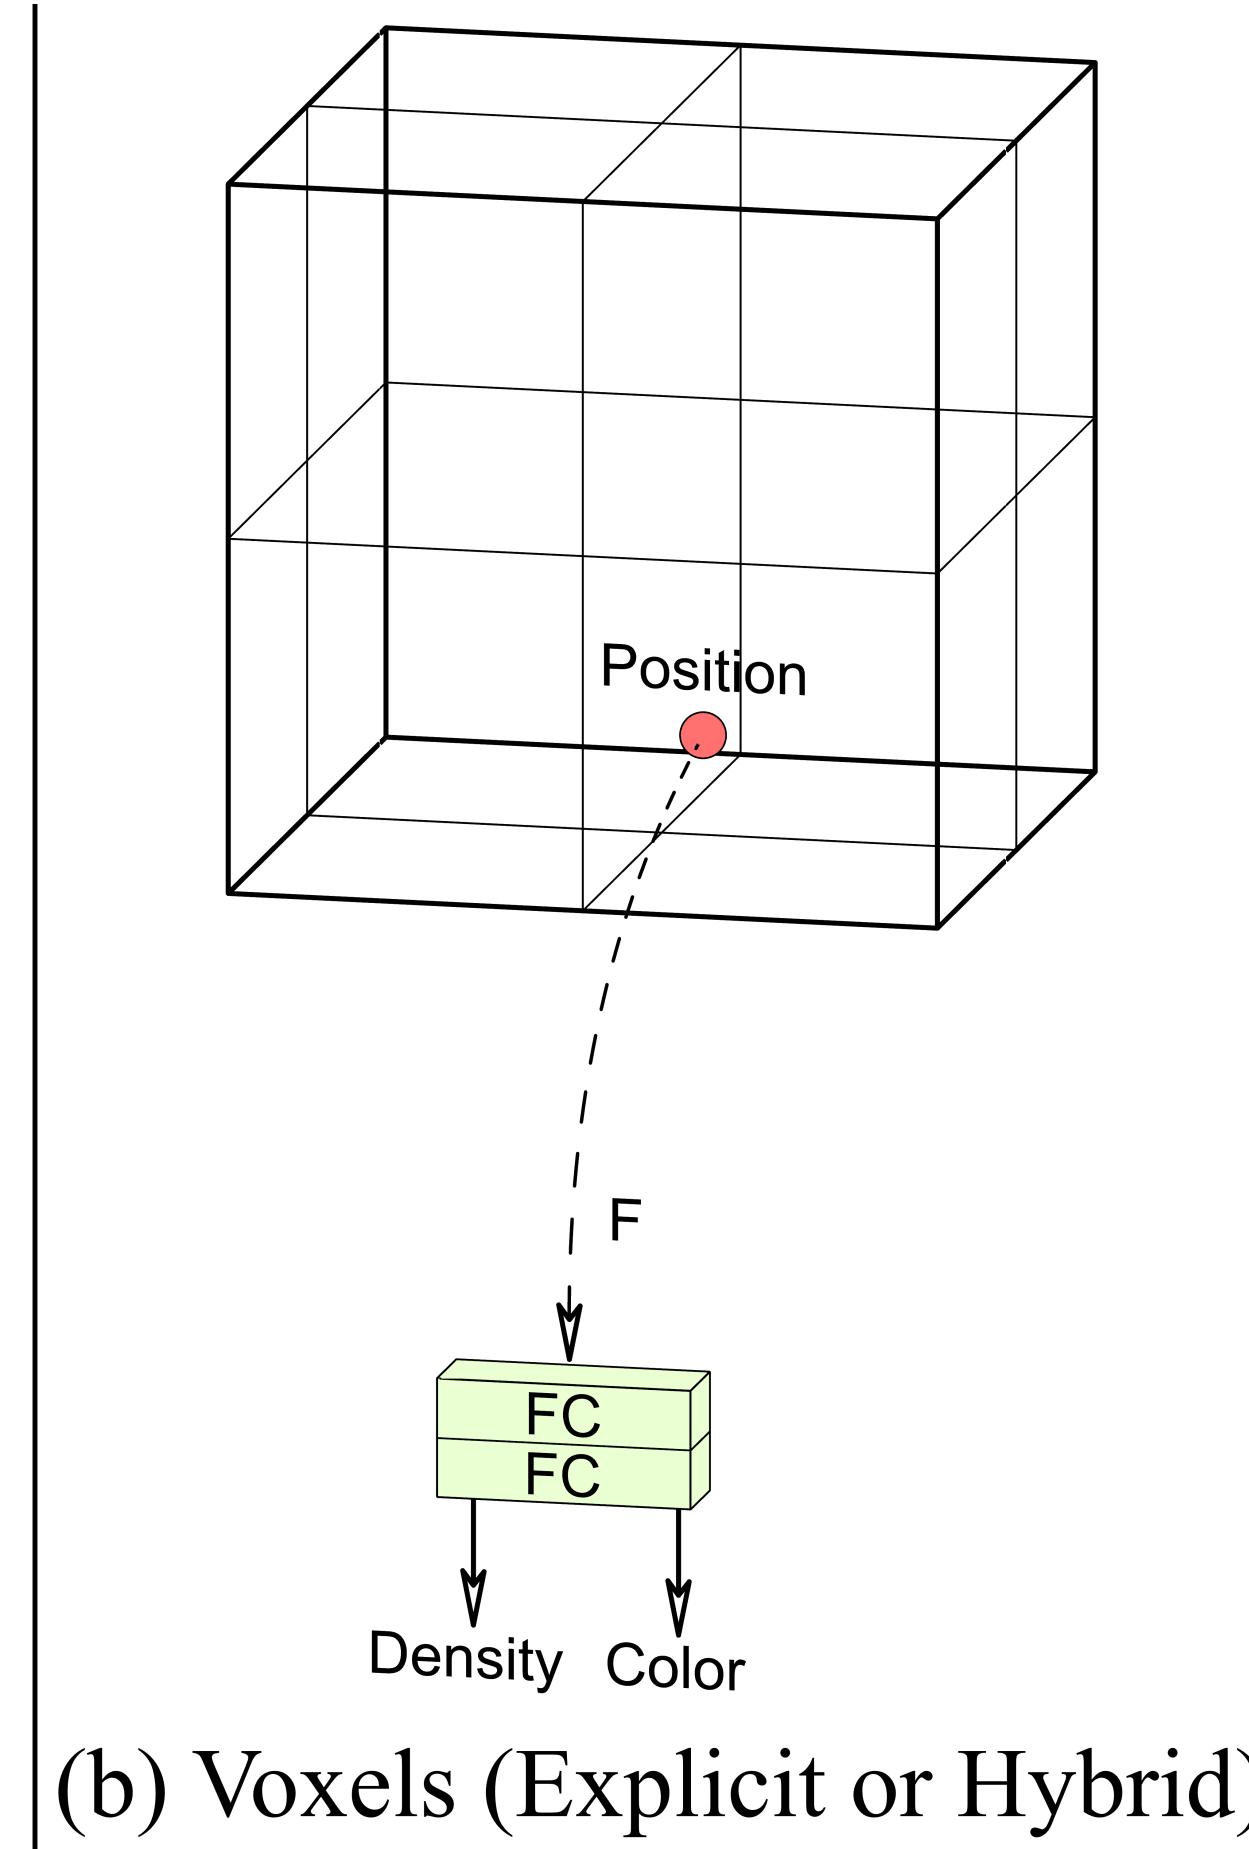
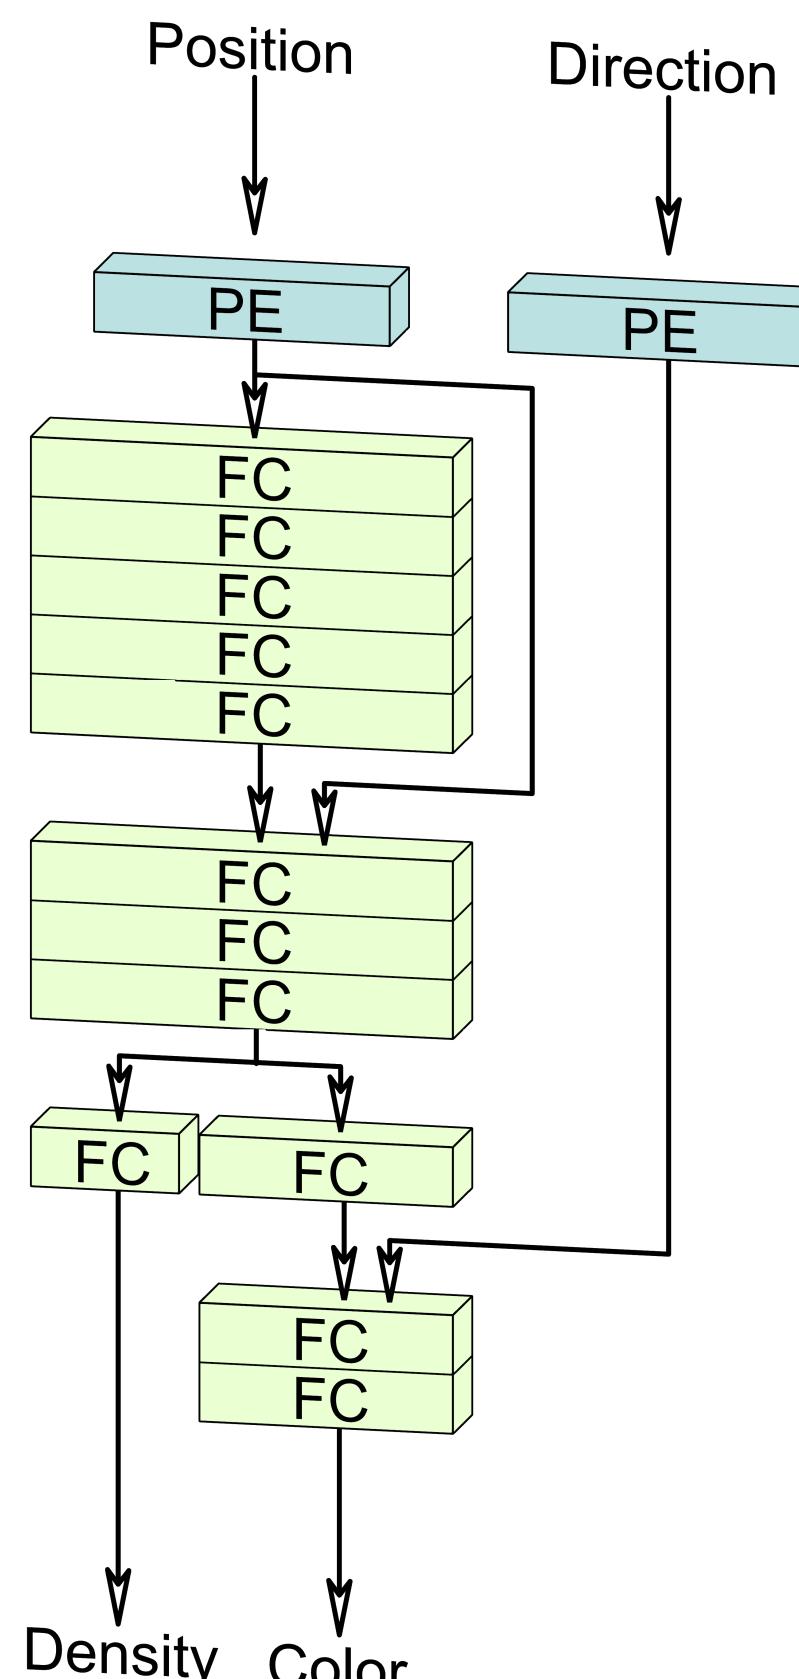
[Efficient Geometry-aware 3D GANs]

# Tri-Planes scene representations



# [Efficient Geometry-aware 3D GANs]

# Tri-Planes scene representations



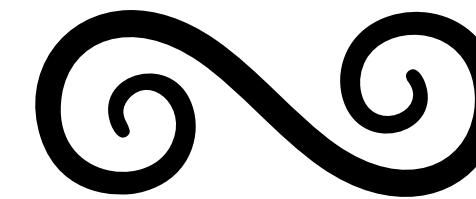
Each of shape  $(T, T, F)$ !

[Efficient Geometry-aware 3D GANs]

- **Tri-Planes are explicit-implicit representations**
- **Tri-Planes  $\cong$  Three planes + tinyMLP + Volume Rendering Equations**
- **Both NeRFs and Tri-Planes are not scalable**

# **Inverse Graphics Problem**

How to model a scene using its captured images?

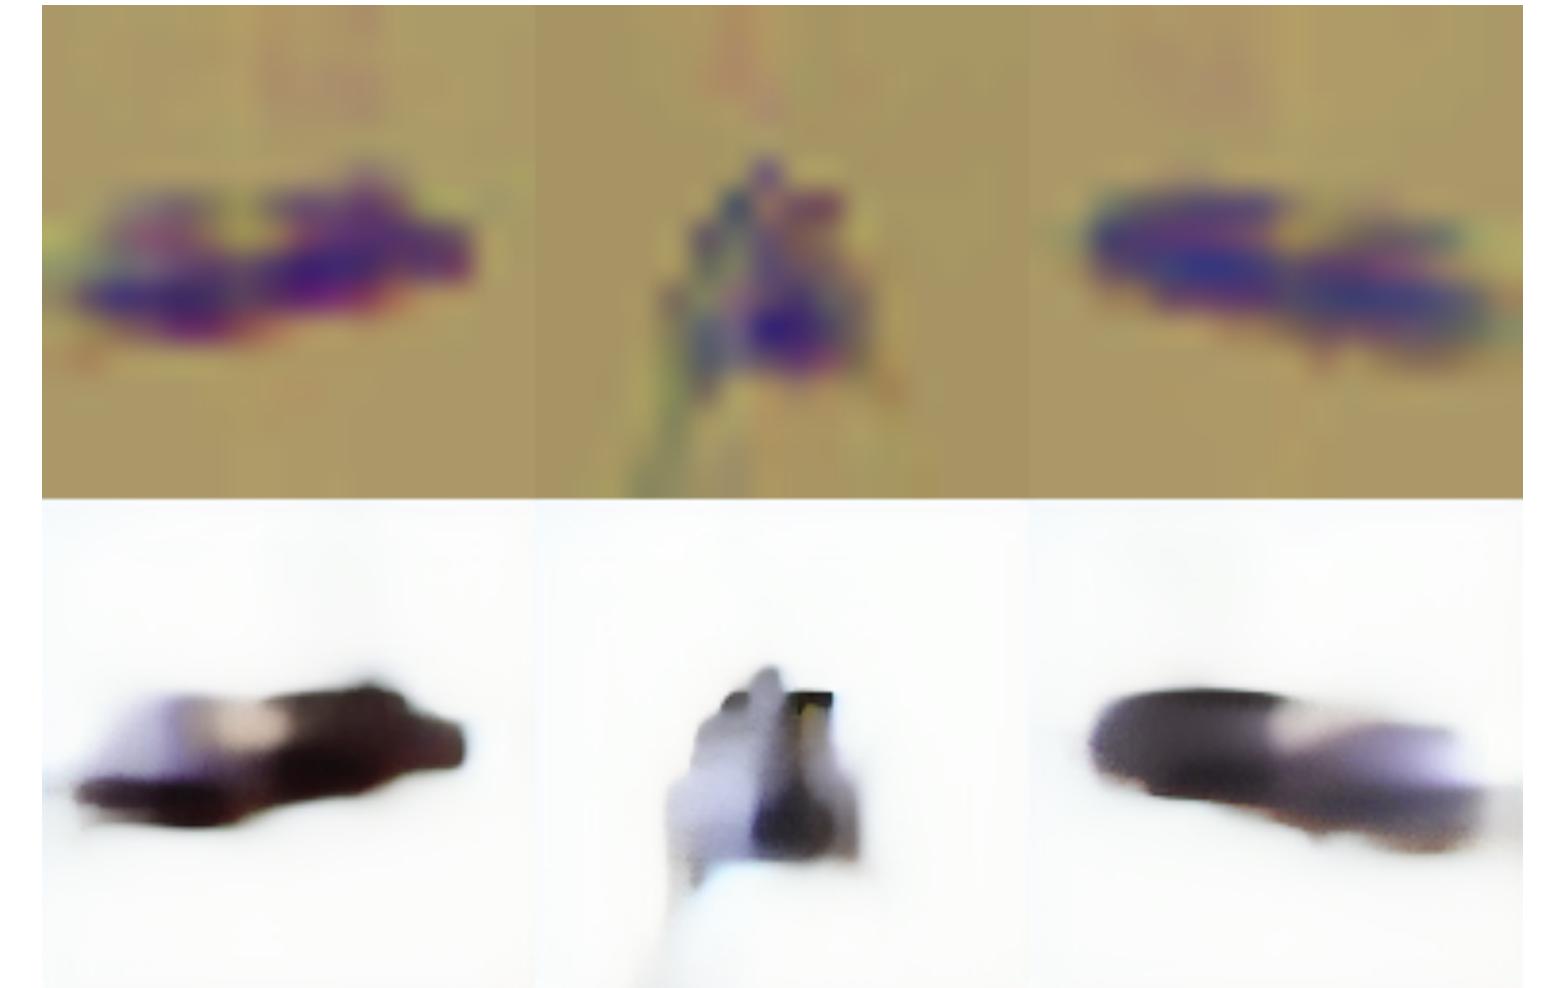


# **(Scaled) Inverse Graphics Problem**

How to model abundantly many scenes at once?

# 3D-aware latent space

- **Goal:** Scale scene representation training in a specially-crafted latent space
  - Improves performances
  - Other applications
- **Neural scene representations main assumption:**
  - The underlying scene behind images is 3D
  - The renderings of the scene are 3D consistent



Tri-Planes trained in a standard AE.

# 3D-aware latent space

- To train scene representations in a latent space, we have to design a 3D-aware latent space

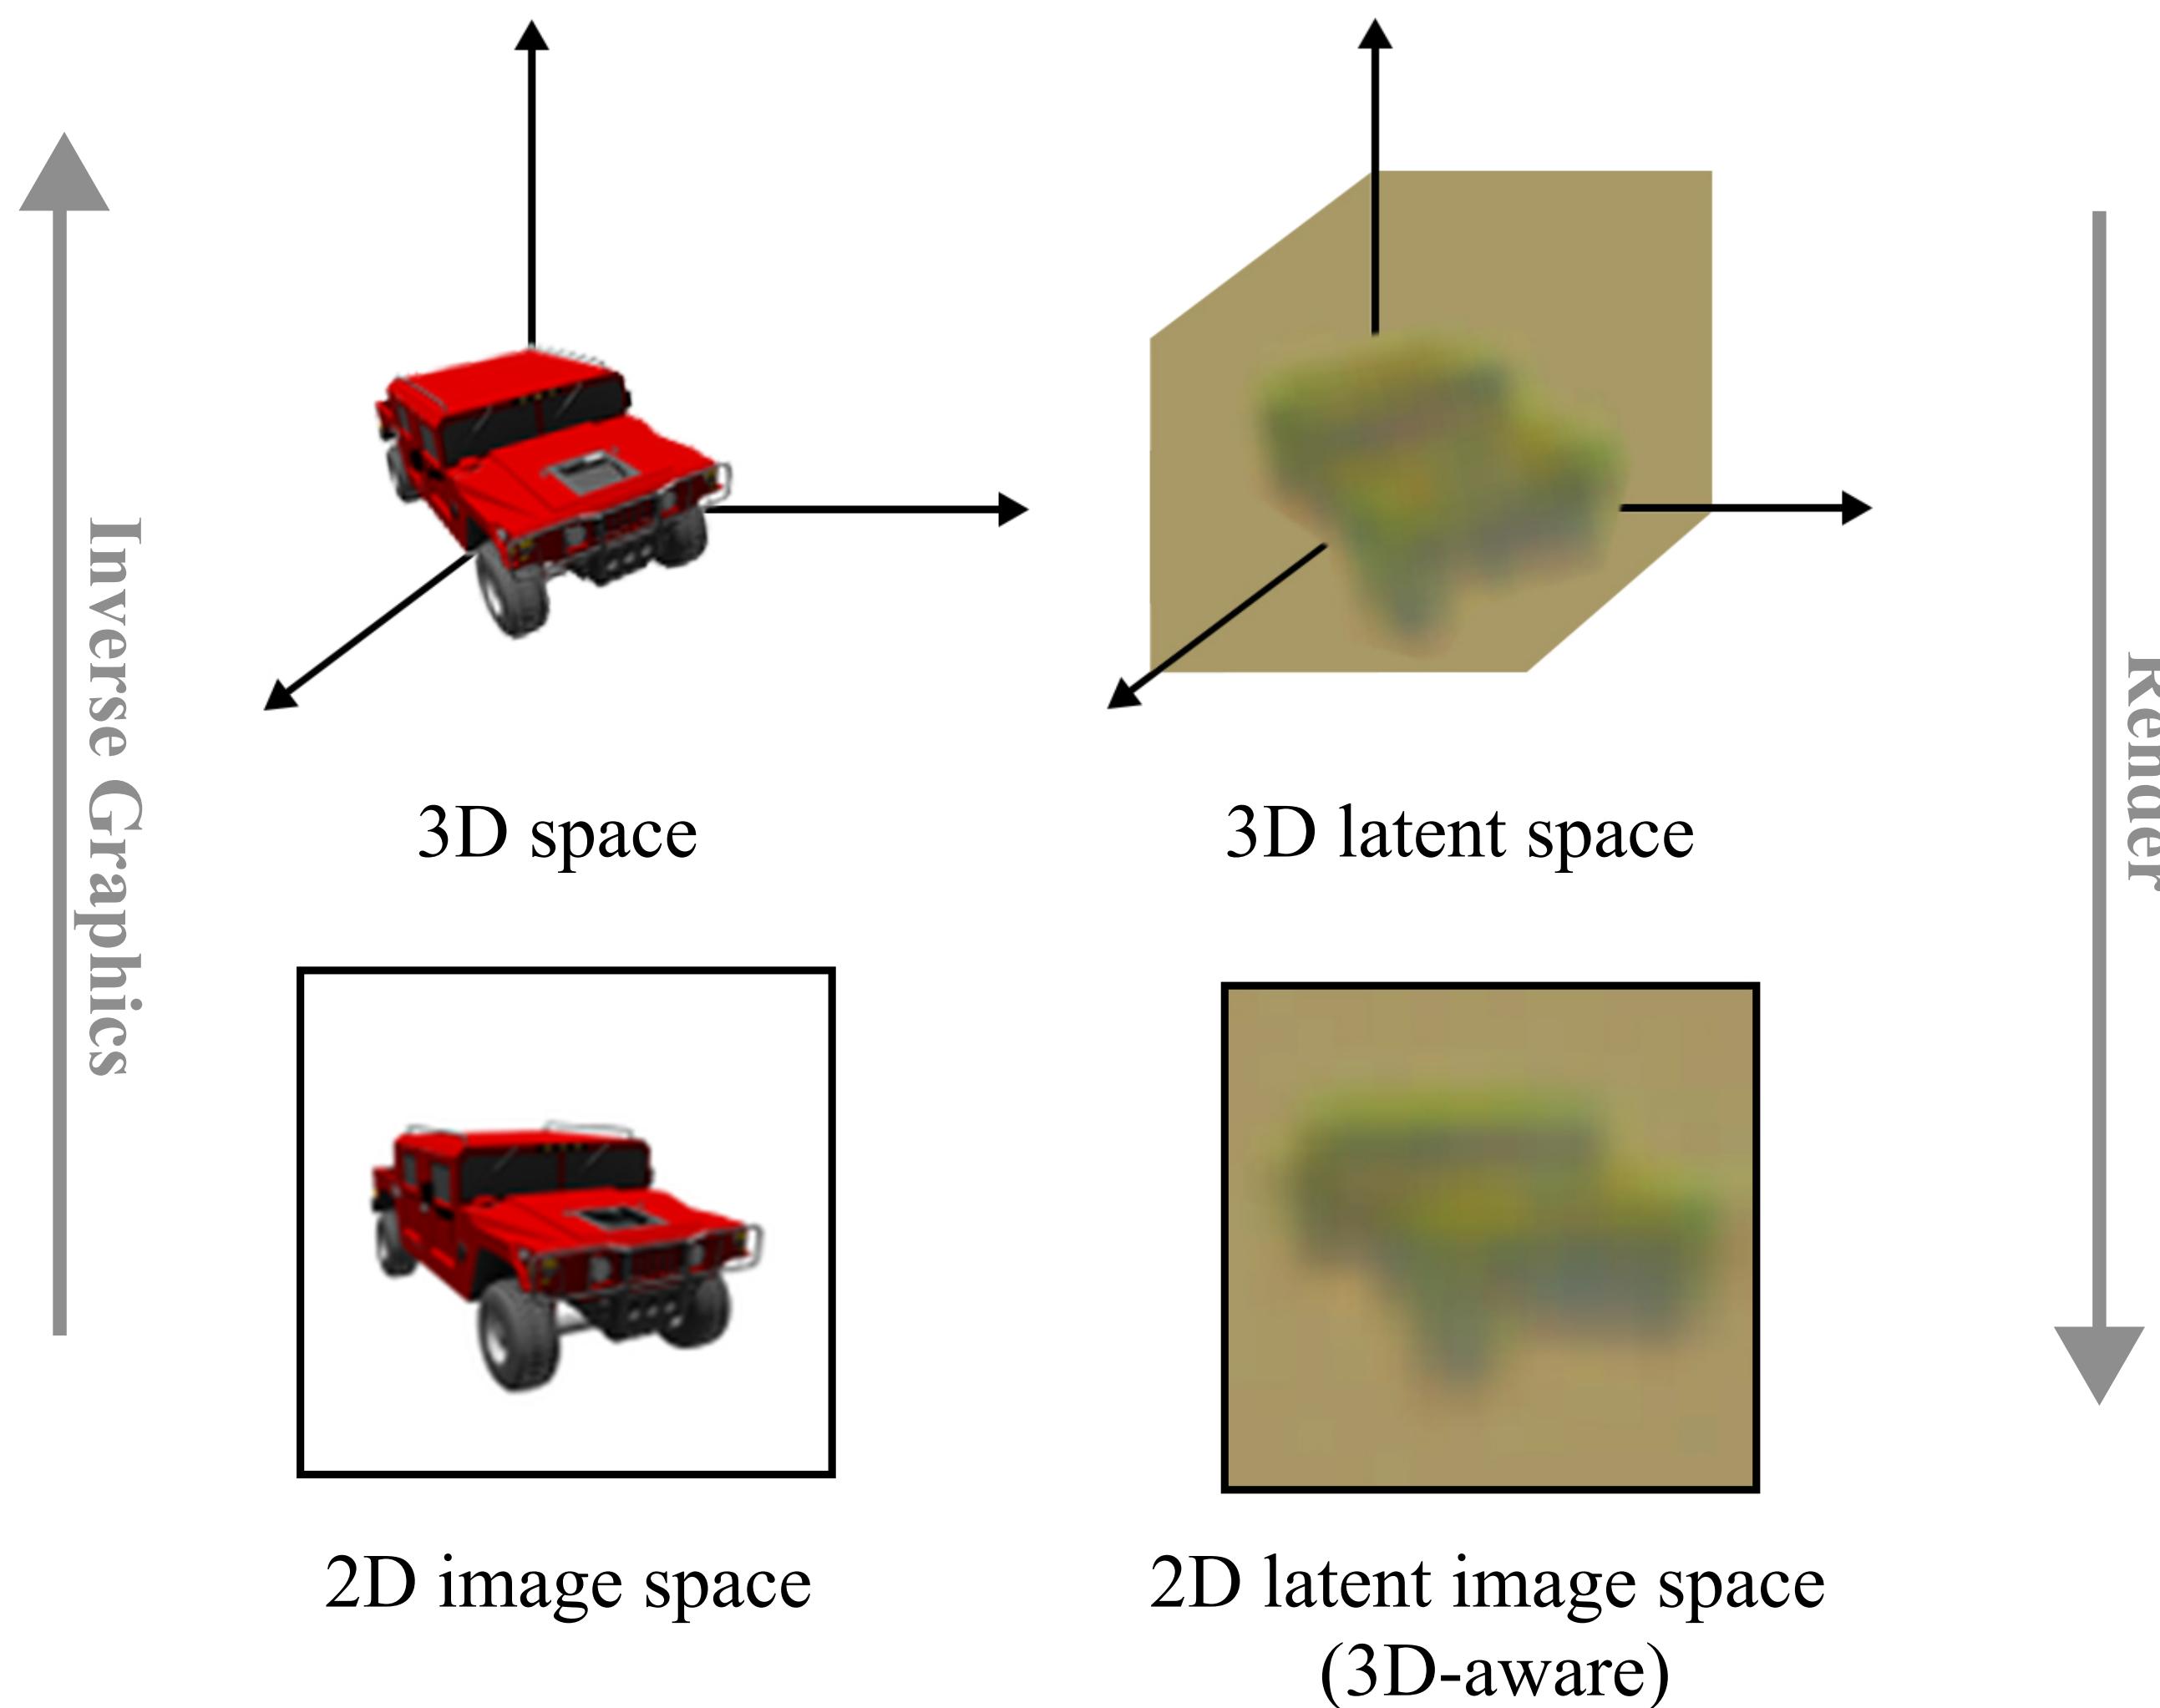
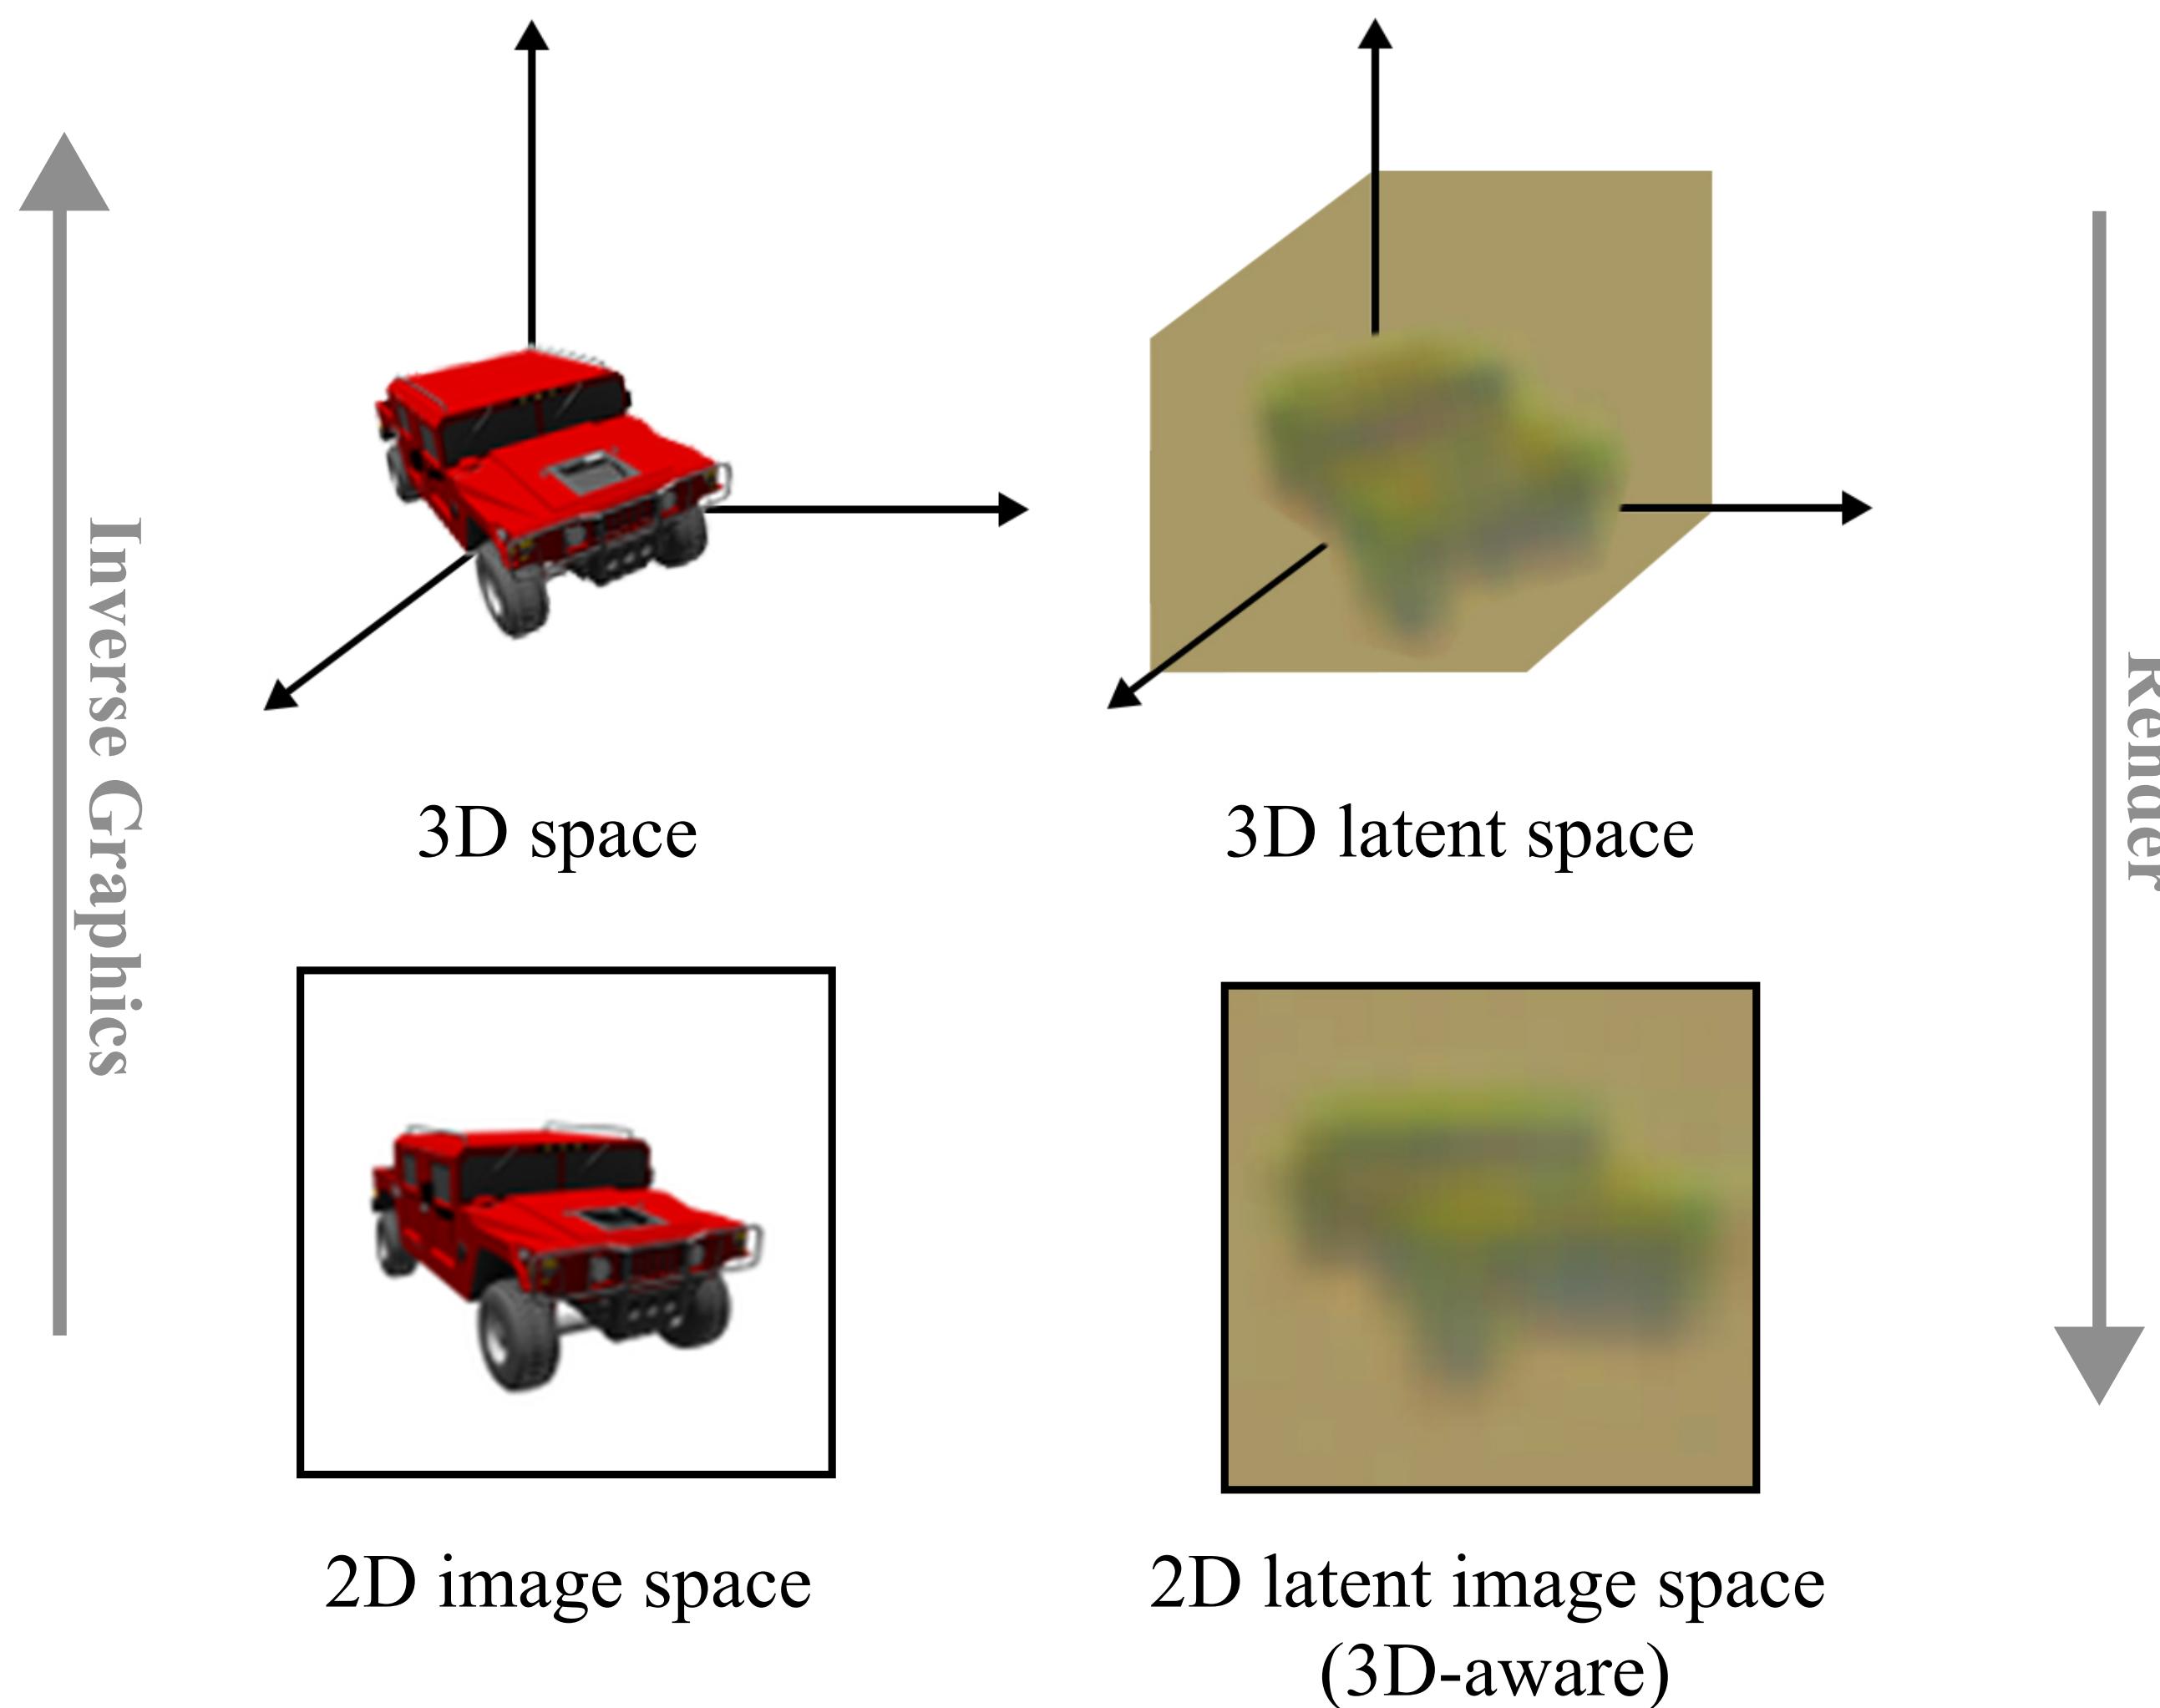


Figure 1. **3D-aware latent space.** We draw inspiration from the relationship between the 3D space and image space and introduce the idea of a 3D latent space. We propose a 3D-aware autoencoder that encodes images into a 3D-aware (2D) latent image space, in which we train our scene representations.

# 3D-aware latent space

- To train scene representations in a latent space, we have to design a 3D-aware latent space



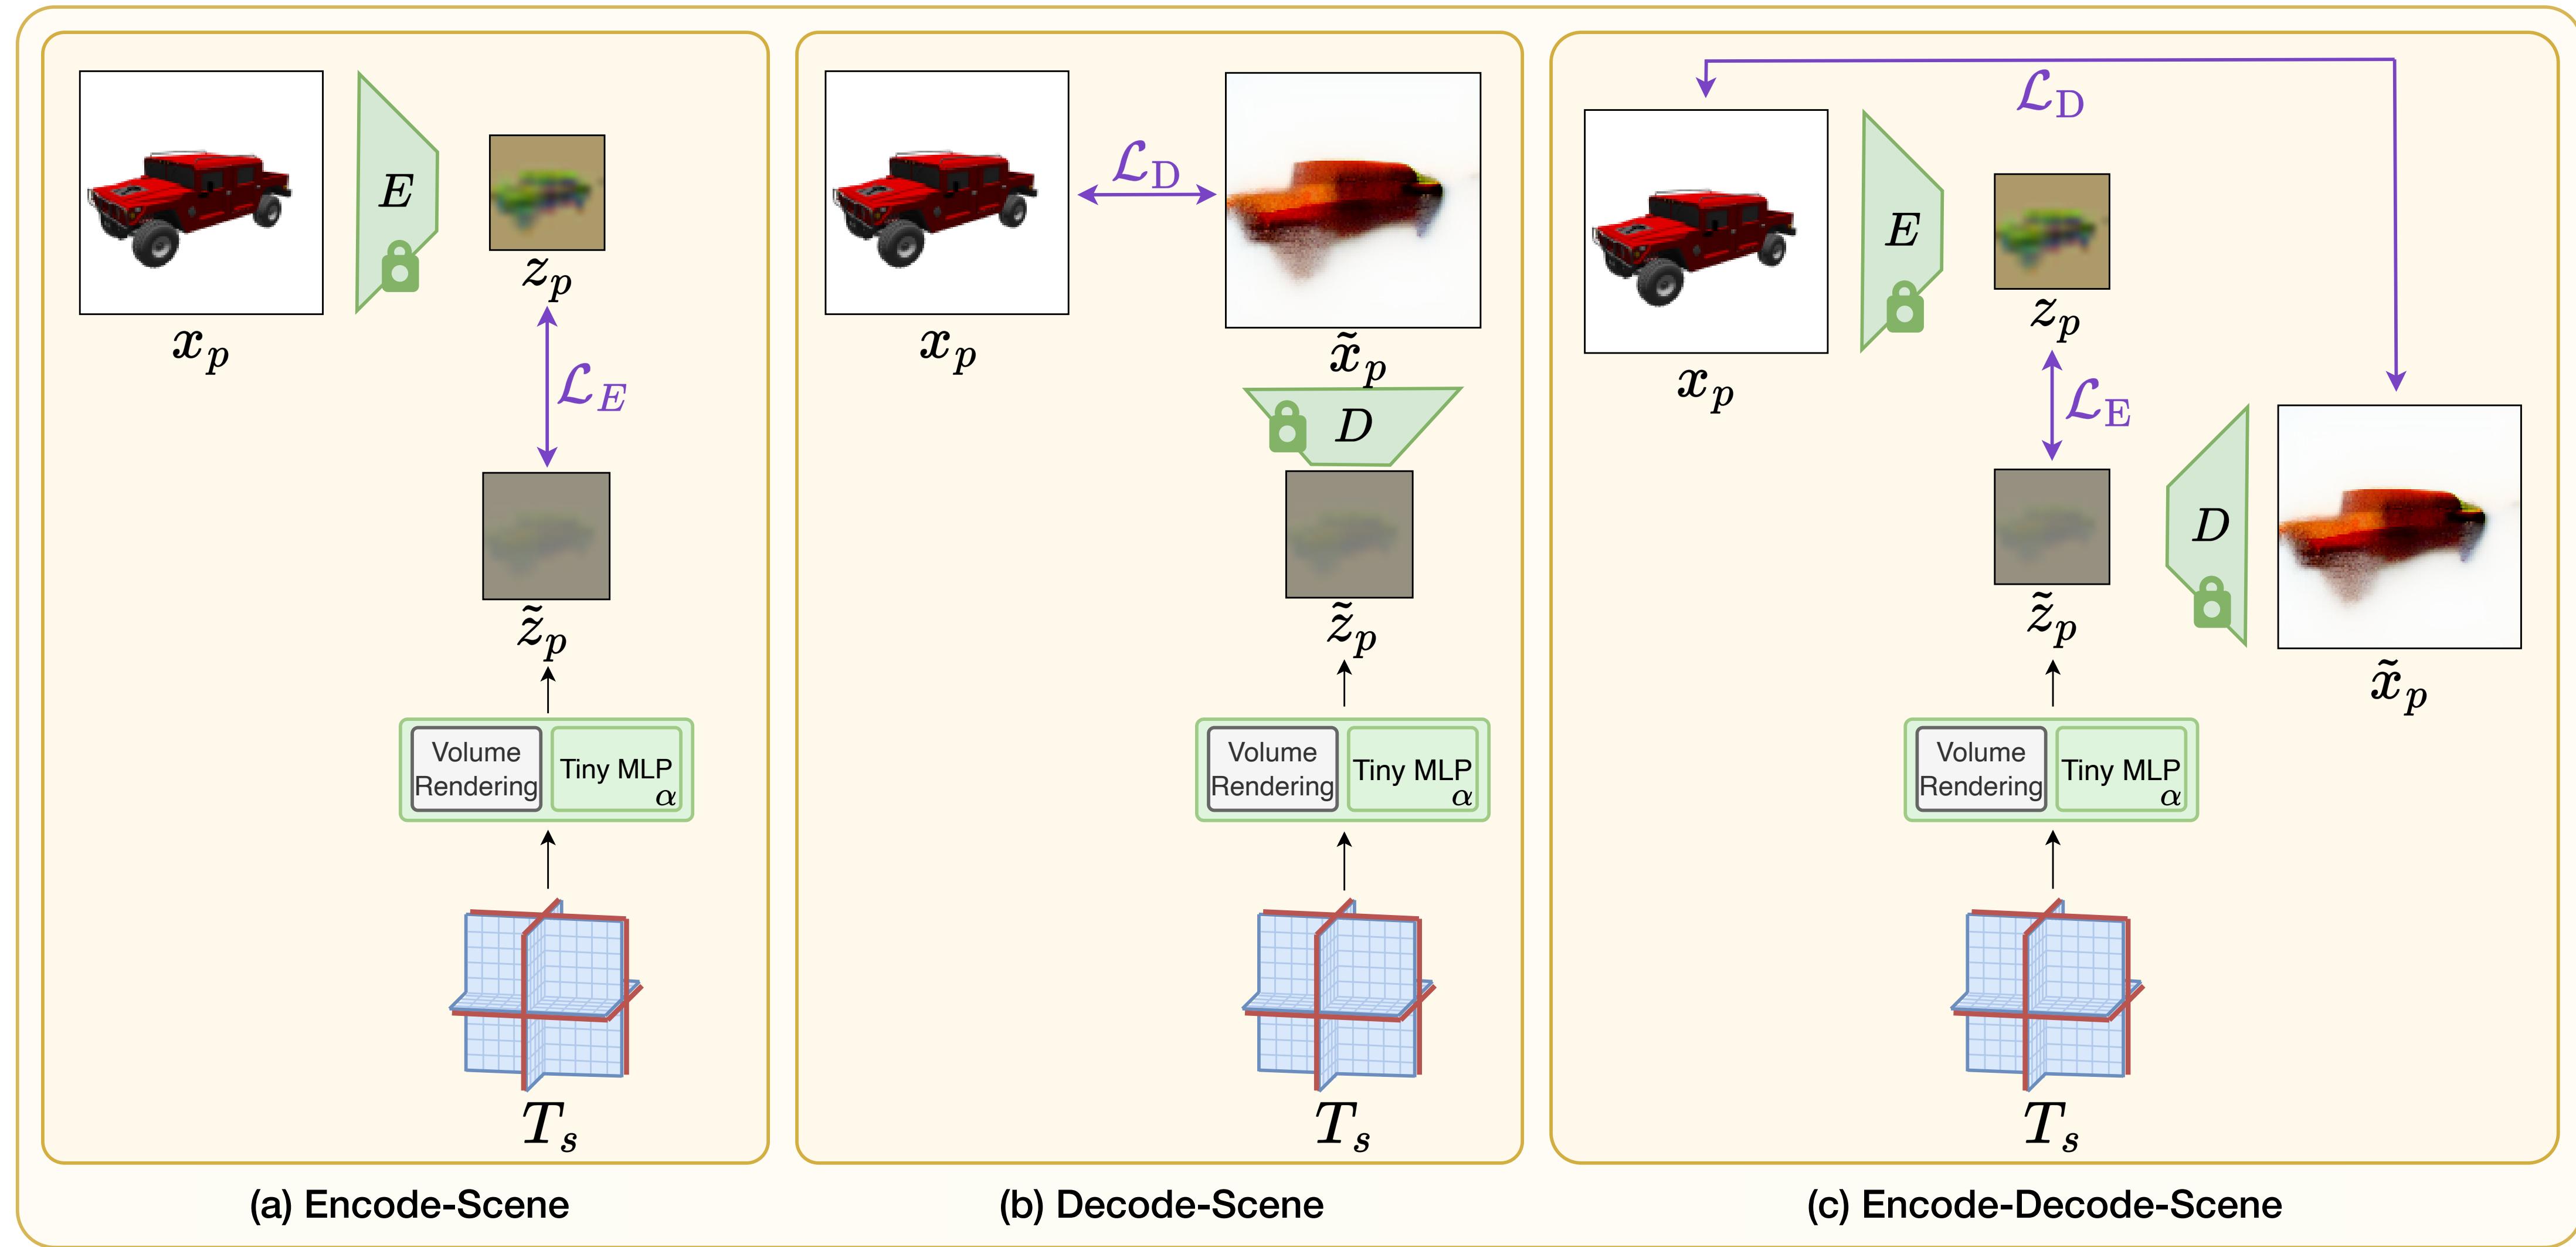
As such, we have two goals:

1. Learn a 3D-aware latent space
2. Leverage that latent space to scale the learning of 3D scenes

Figure 1. **3D-aware latent space.** We draw inspiration from the relationship between the 3D space and image space and introduce the idea of a 3D latent space. We propose a 3D-aware autoencoder that encodes images into a 3D-aware (2D) latent image space, in which we train our scene representations.

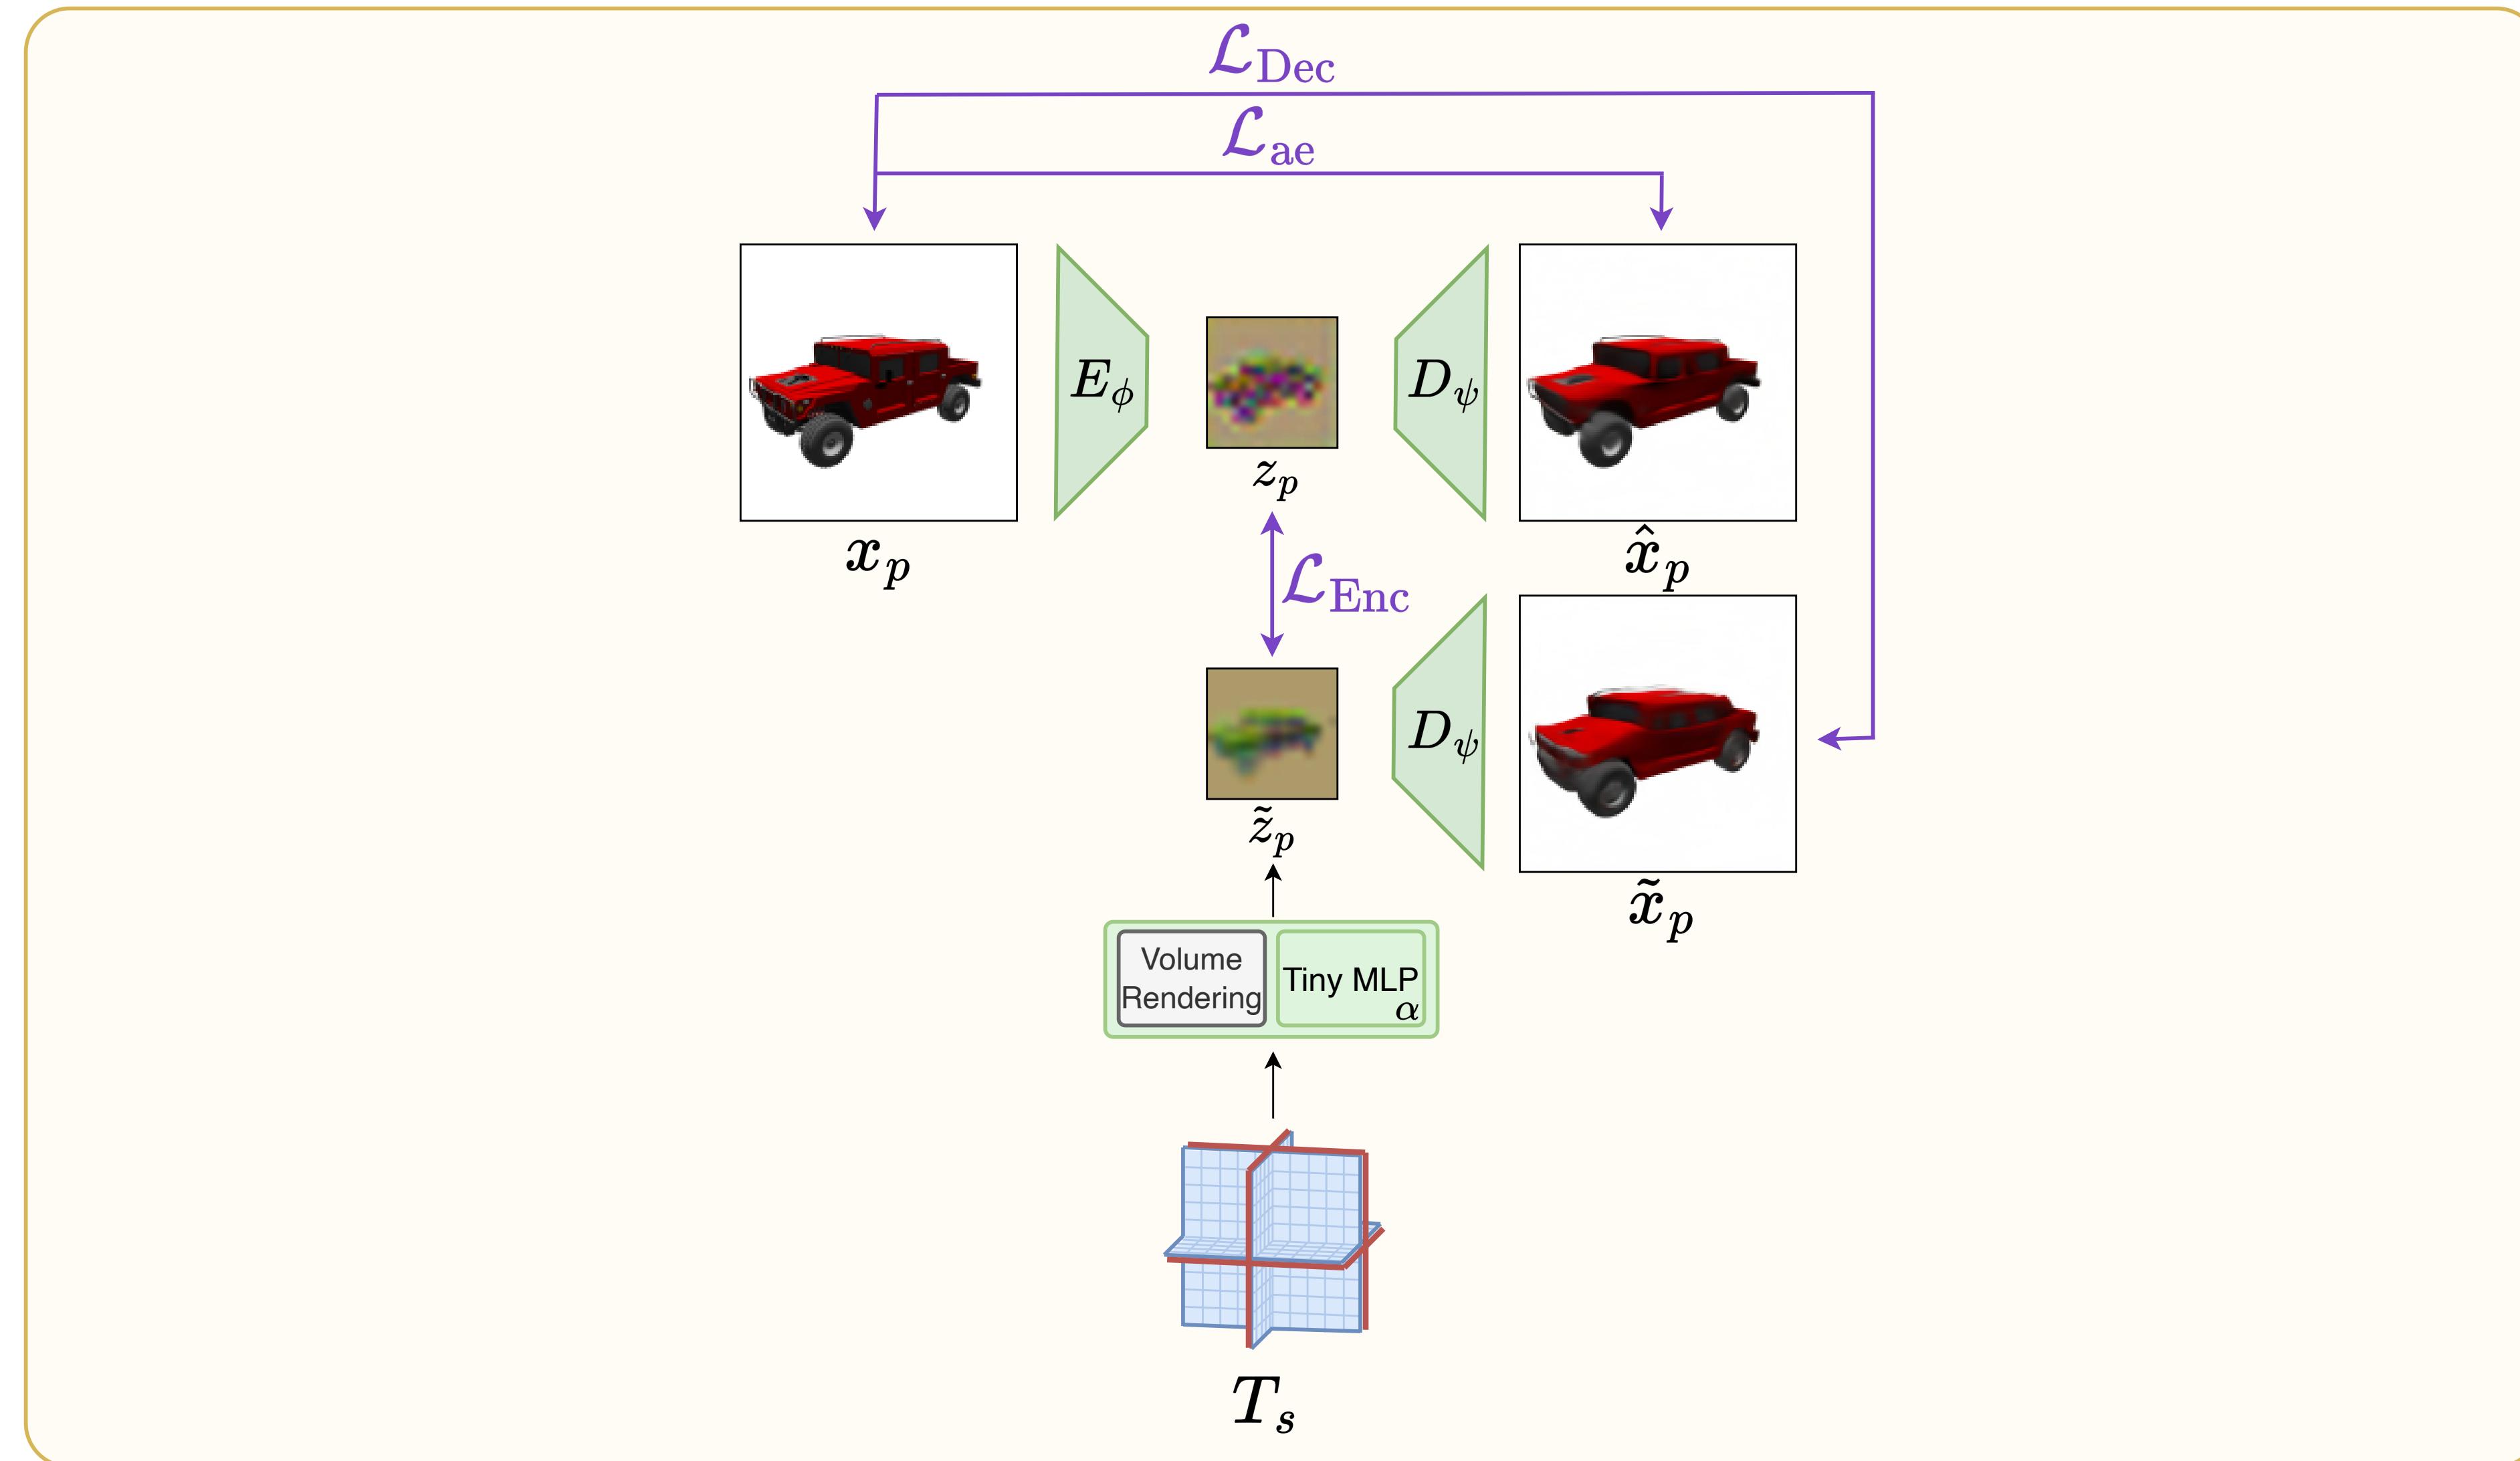
**How to learn latent scenes given a 3D-aware latent space?**

# How to learn latent scenes given a 3D-aware latent space?

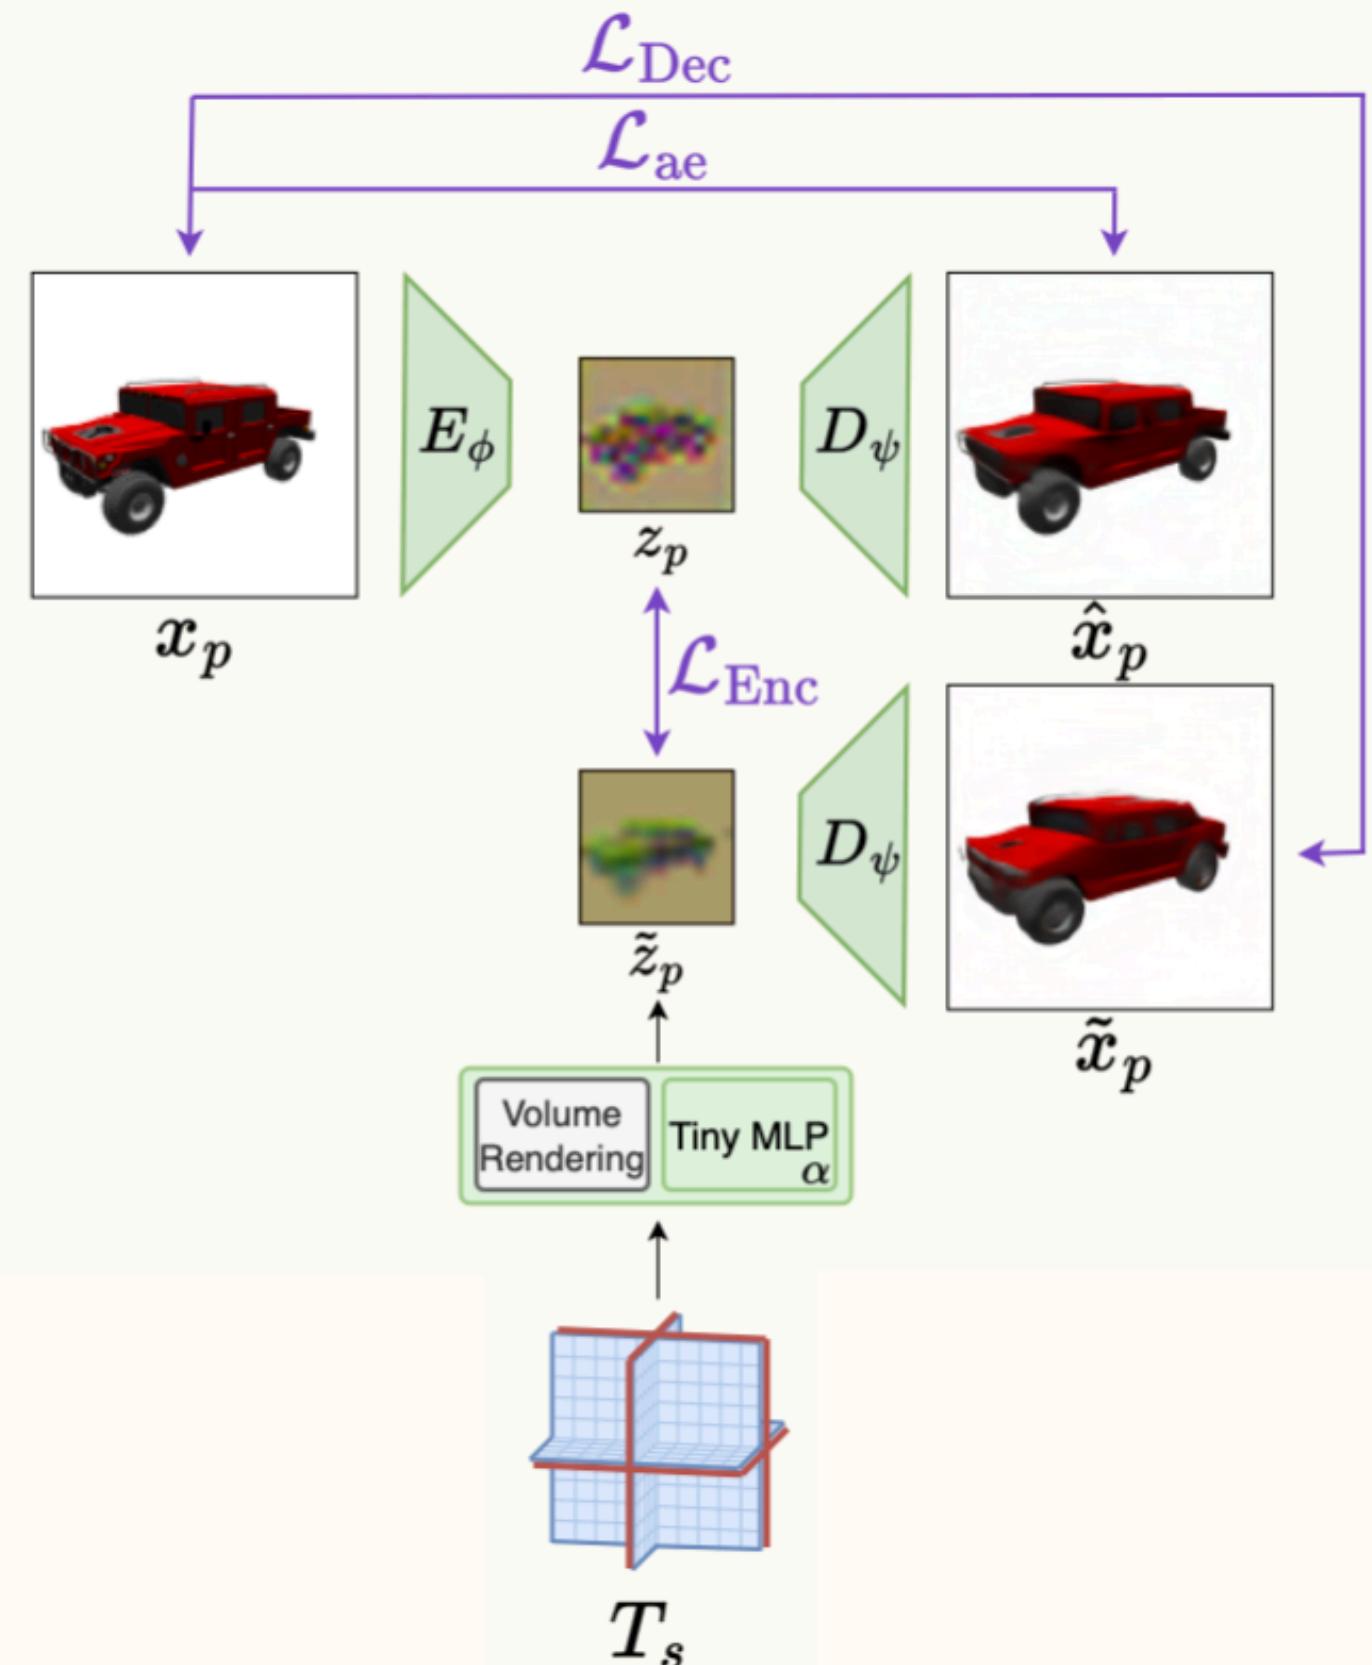


# **How to learn a 3D-aware latent space?**

# How to learn a 3D-aware latent space?



# How to learn a 3D-aware latent space?

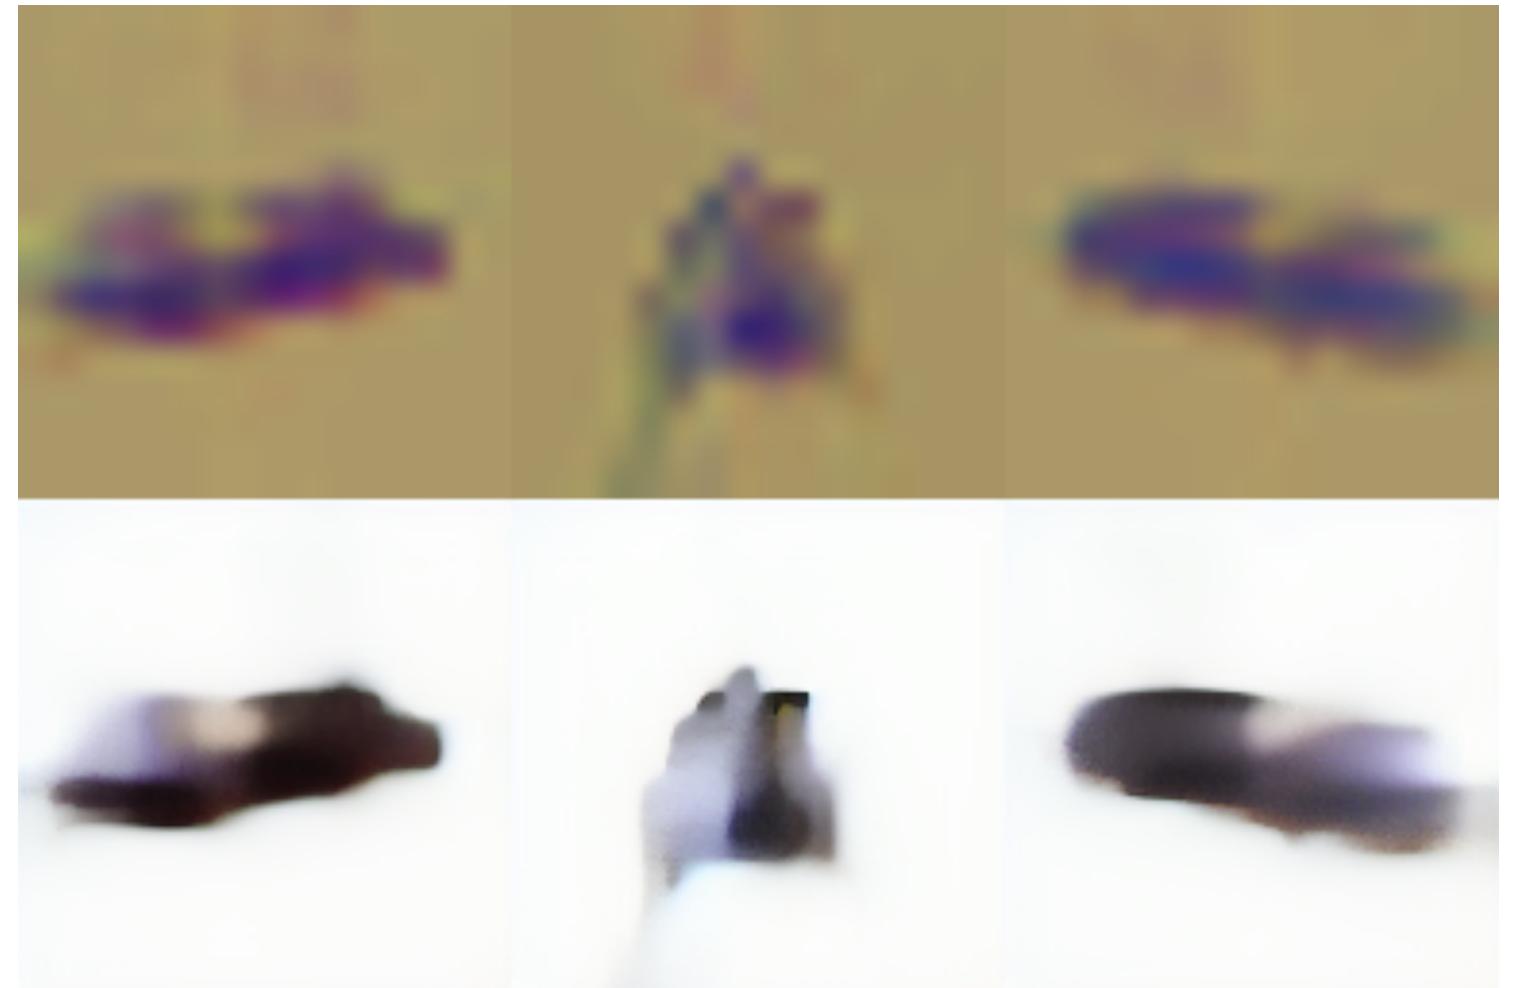


$$\min_{\phi, \psi, \alpha, T} \lambda_{\text{ae}} \mathcal{L}_{\text{ae}}(\phi, \psi) + \lambda_{\text{Enc}} \mathcal{L}_{\text{Enc}}(\phi, \alpha, T) + \lambda_{\text{Dec}} \mathcal{L}_{\text{Dec}}(\psi, \alpha, T) ,$$

with

$$\left\{ \begin{array}{l} \mathcal{L}_{\text{ae}}(\phi, \psi) = \mathbb{E}_{x_p} \|x_p - D_\psi(E_\phi(x_p))\| , \\ \mathcal{L}_{\text{Enc}}(\phi, \alpha, T) = \mathbb{E}_{x_p} \|E_\phi(x_p) - \mathcal{R}_\alpha(T, p)\| , \\ \mathcal{L}_{\text{Dec}}(\psi, \alpha, T) = \mathbb{E}_{x_p} \|x_p - D_\psi(\mathcal{R}_\alpha(T, p))\| , \end{array} \right.$$

# 3D-aware latent space

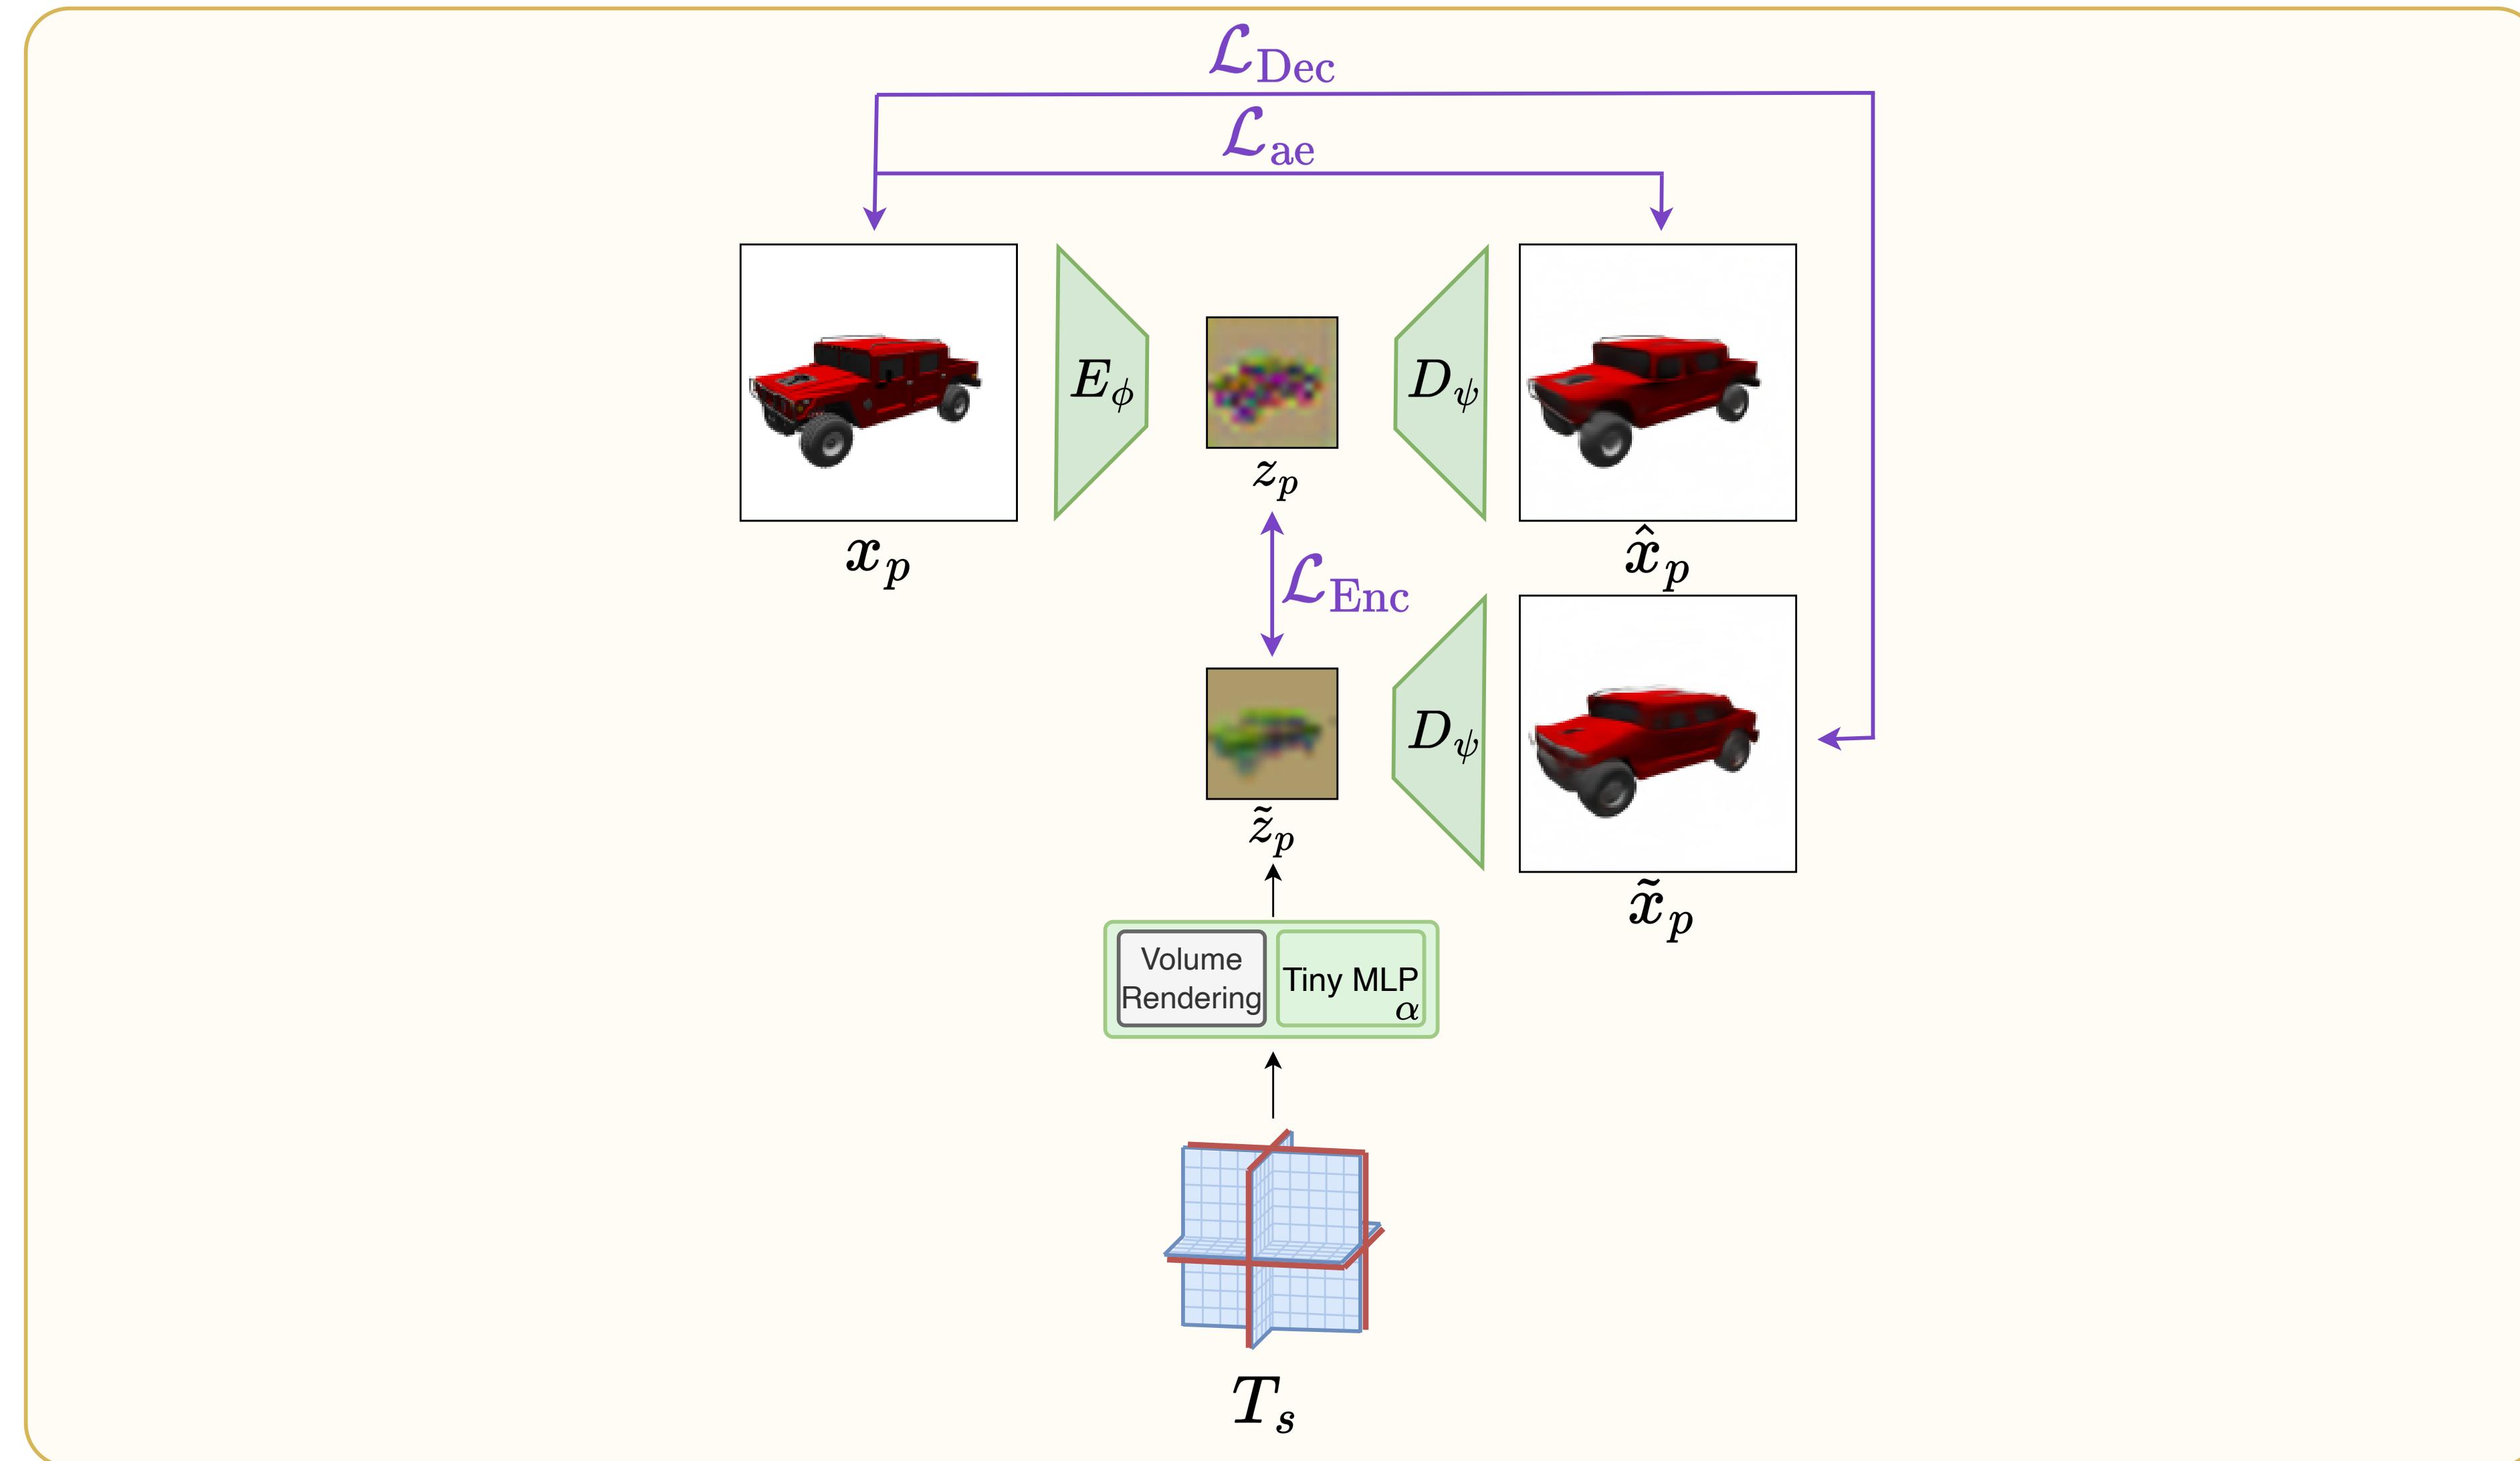


Tri-Planes trained in a standard AE.

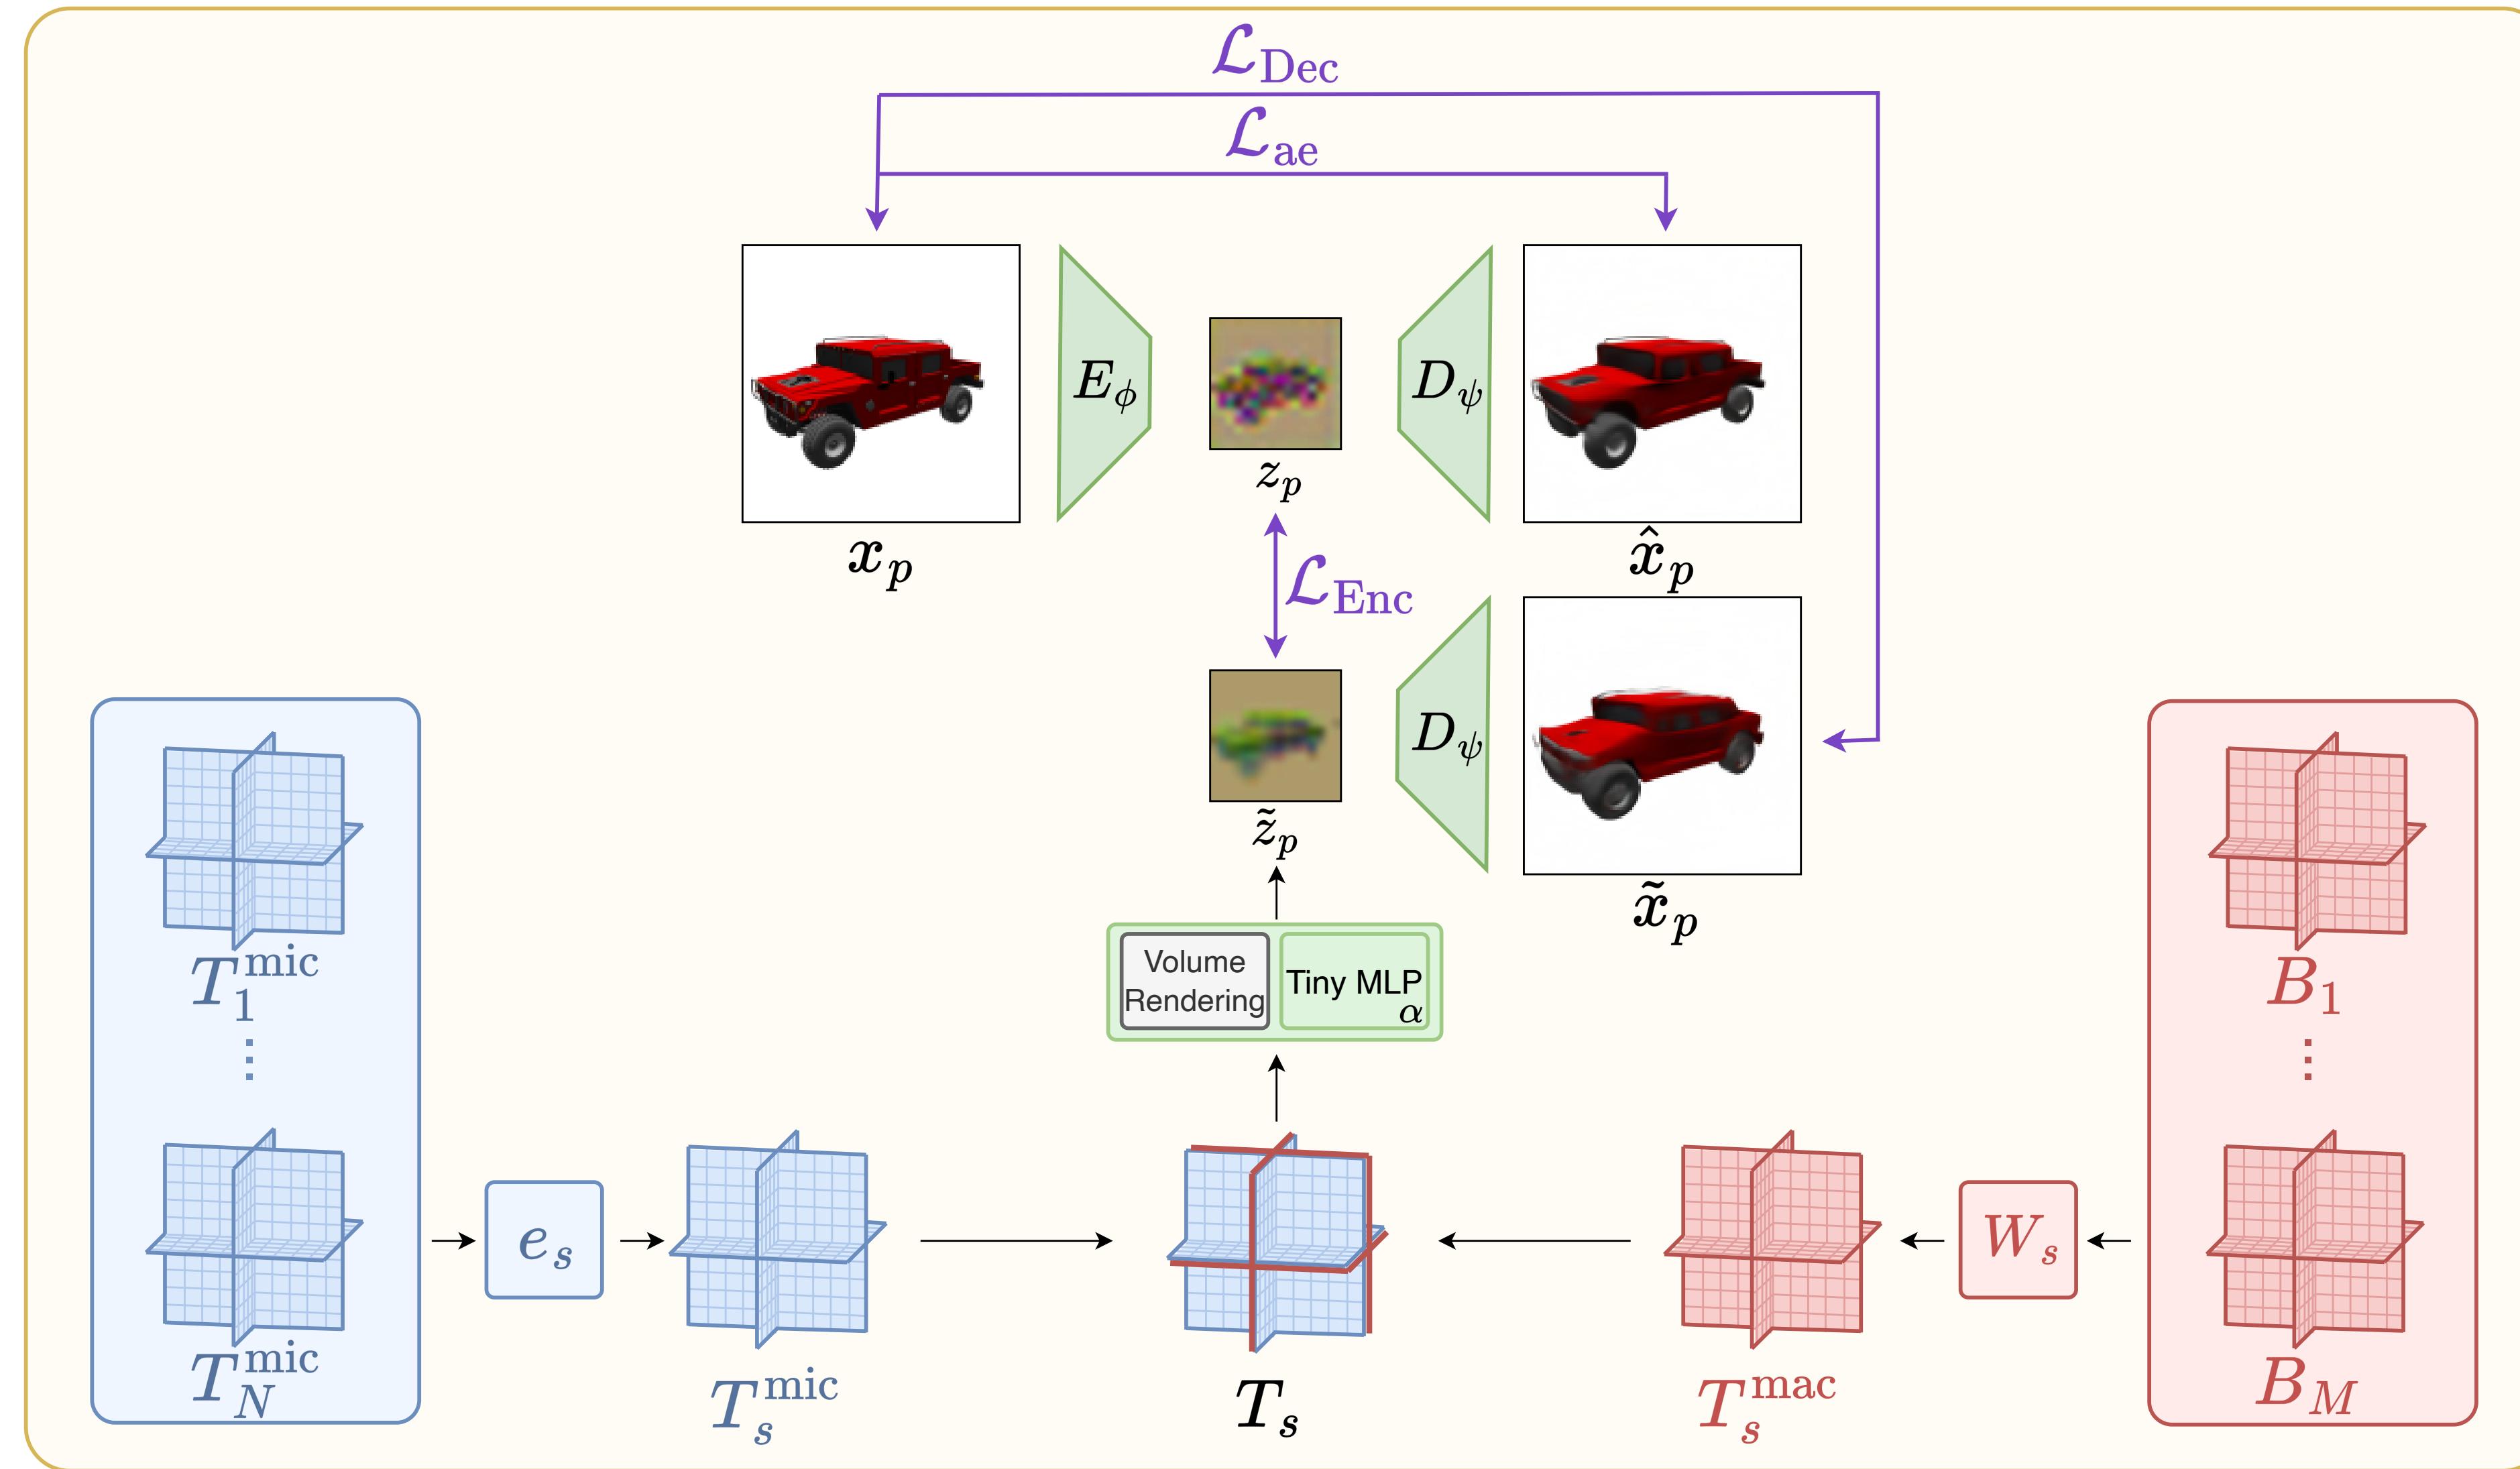


Tri-Planes trained in a 3Da-AE.

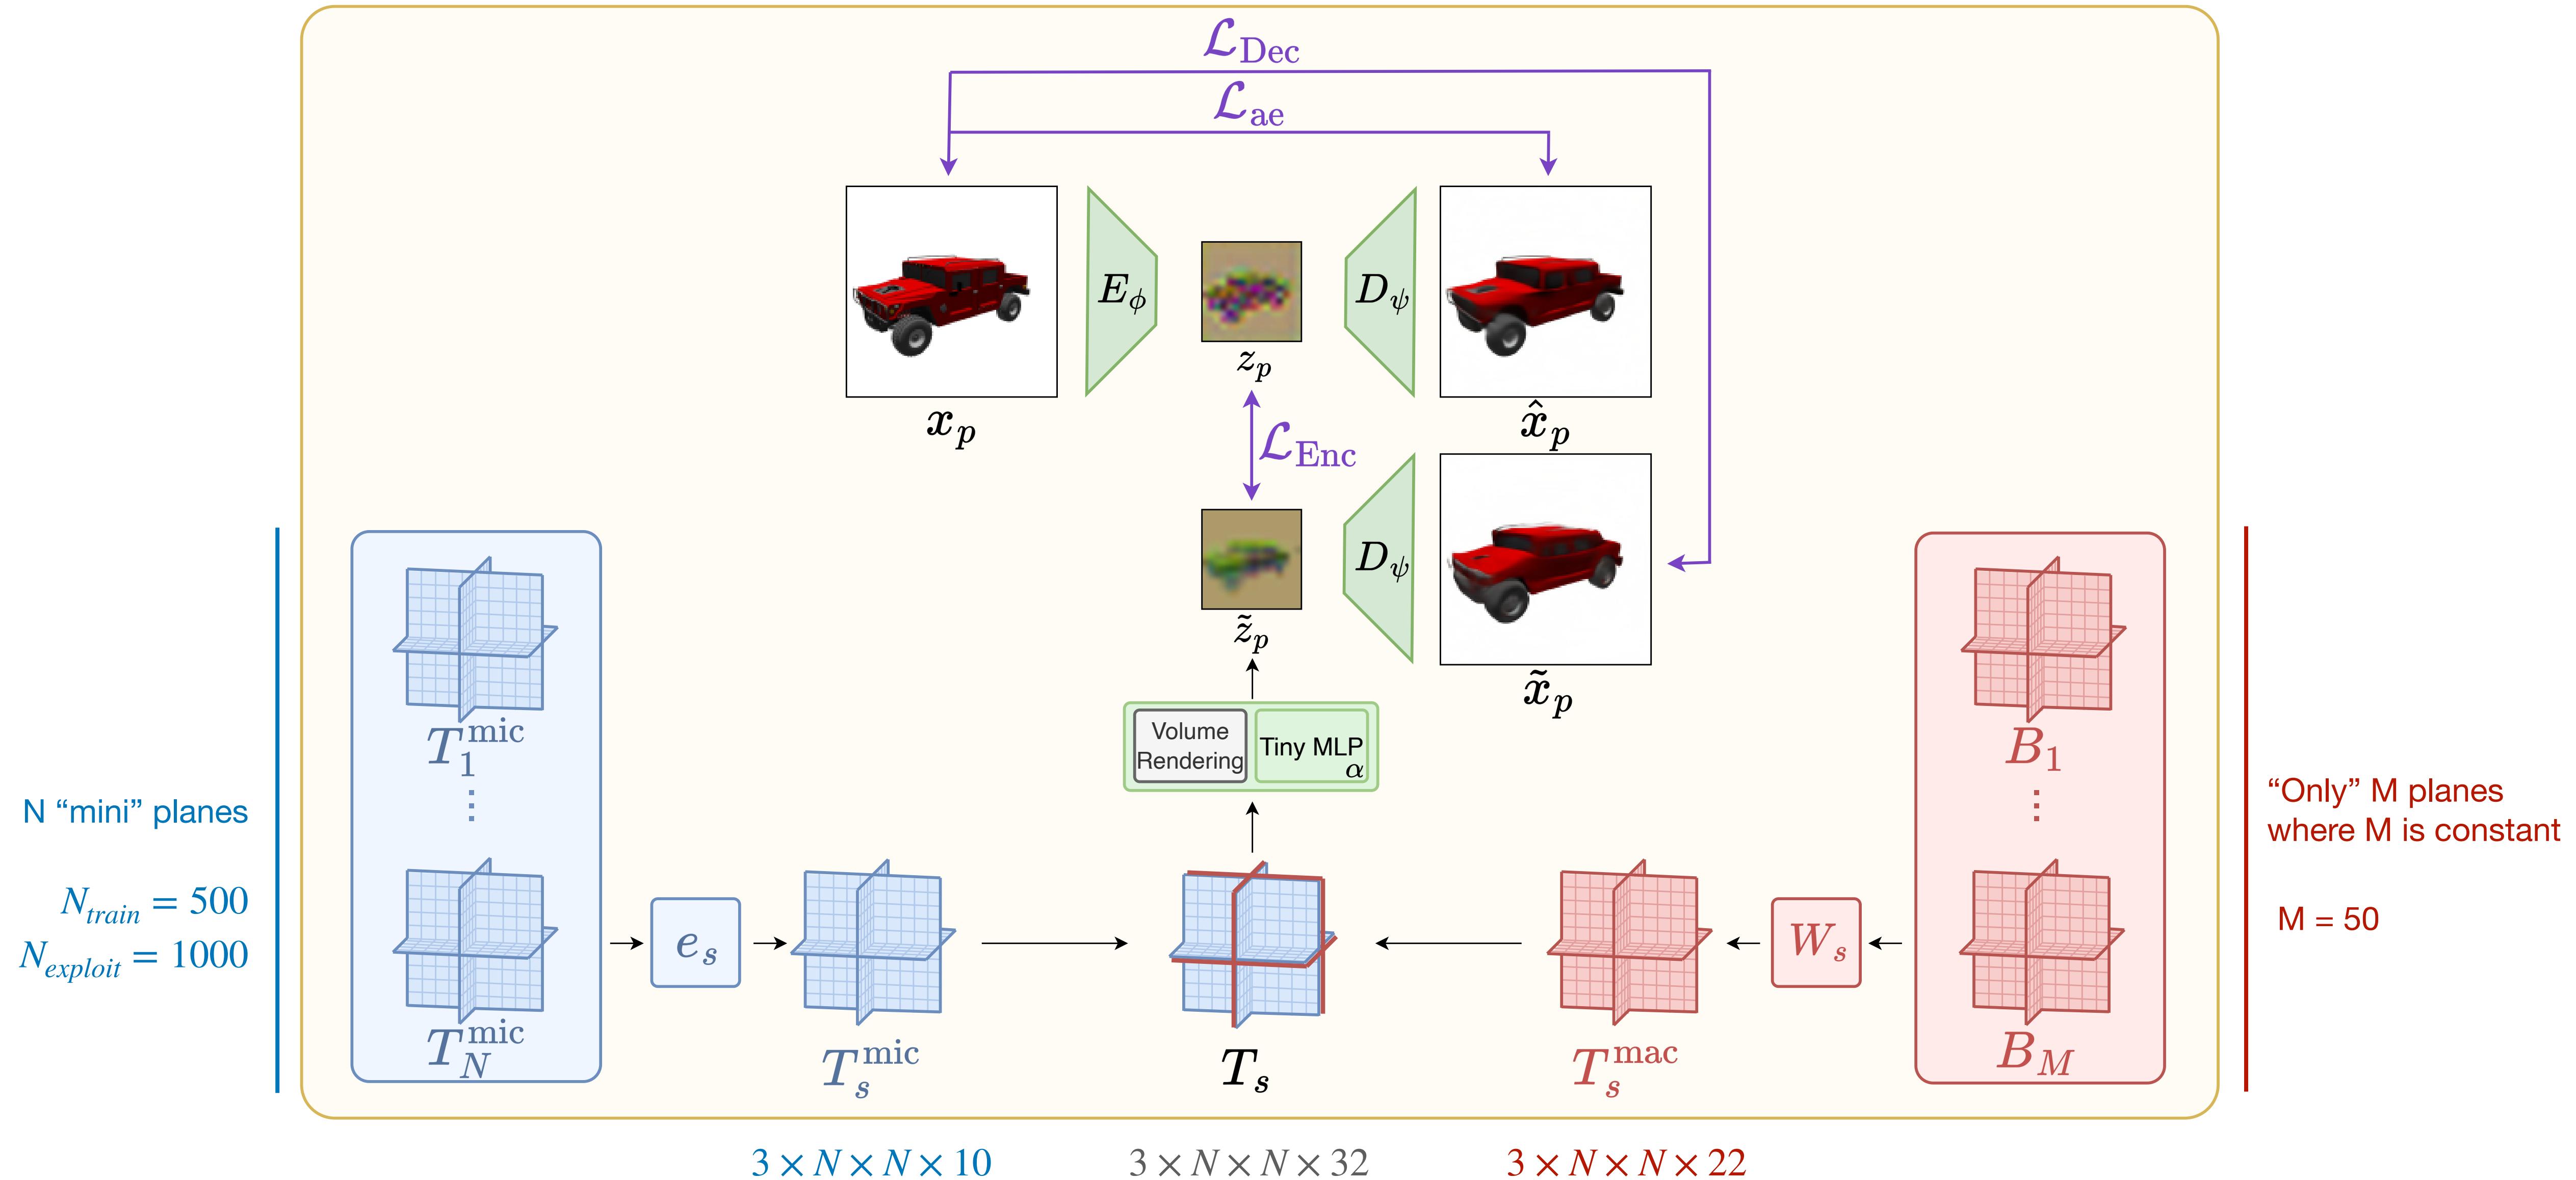
# How to learn a 3D-aware latent space?



# How to further scale training in a 3D-aware latent space?

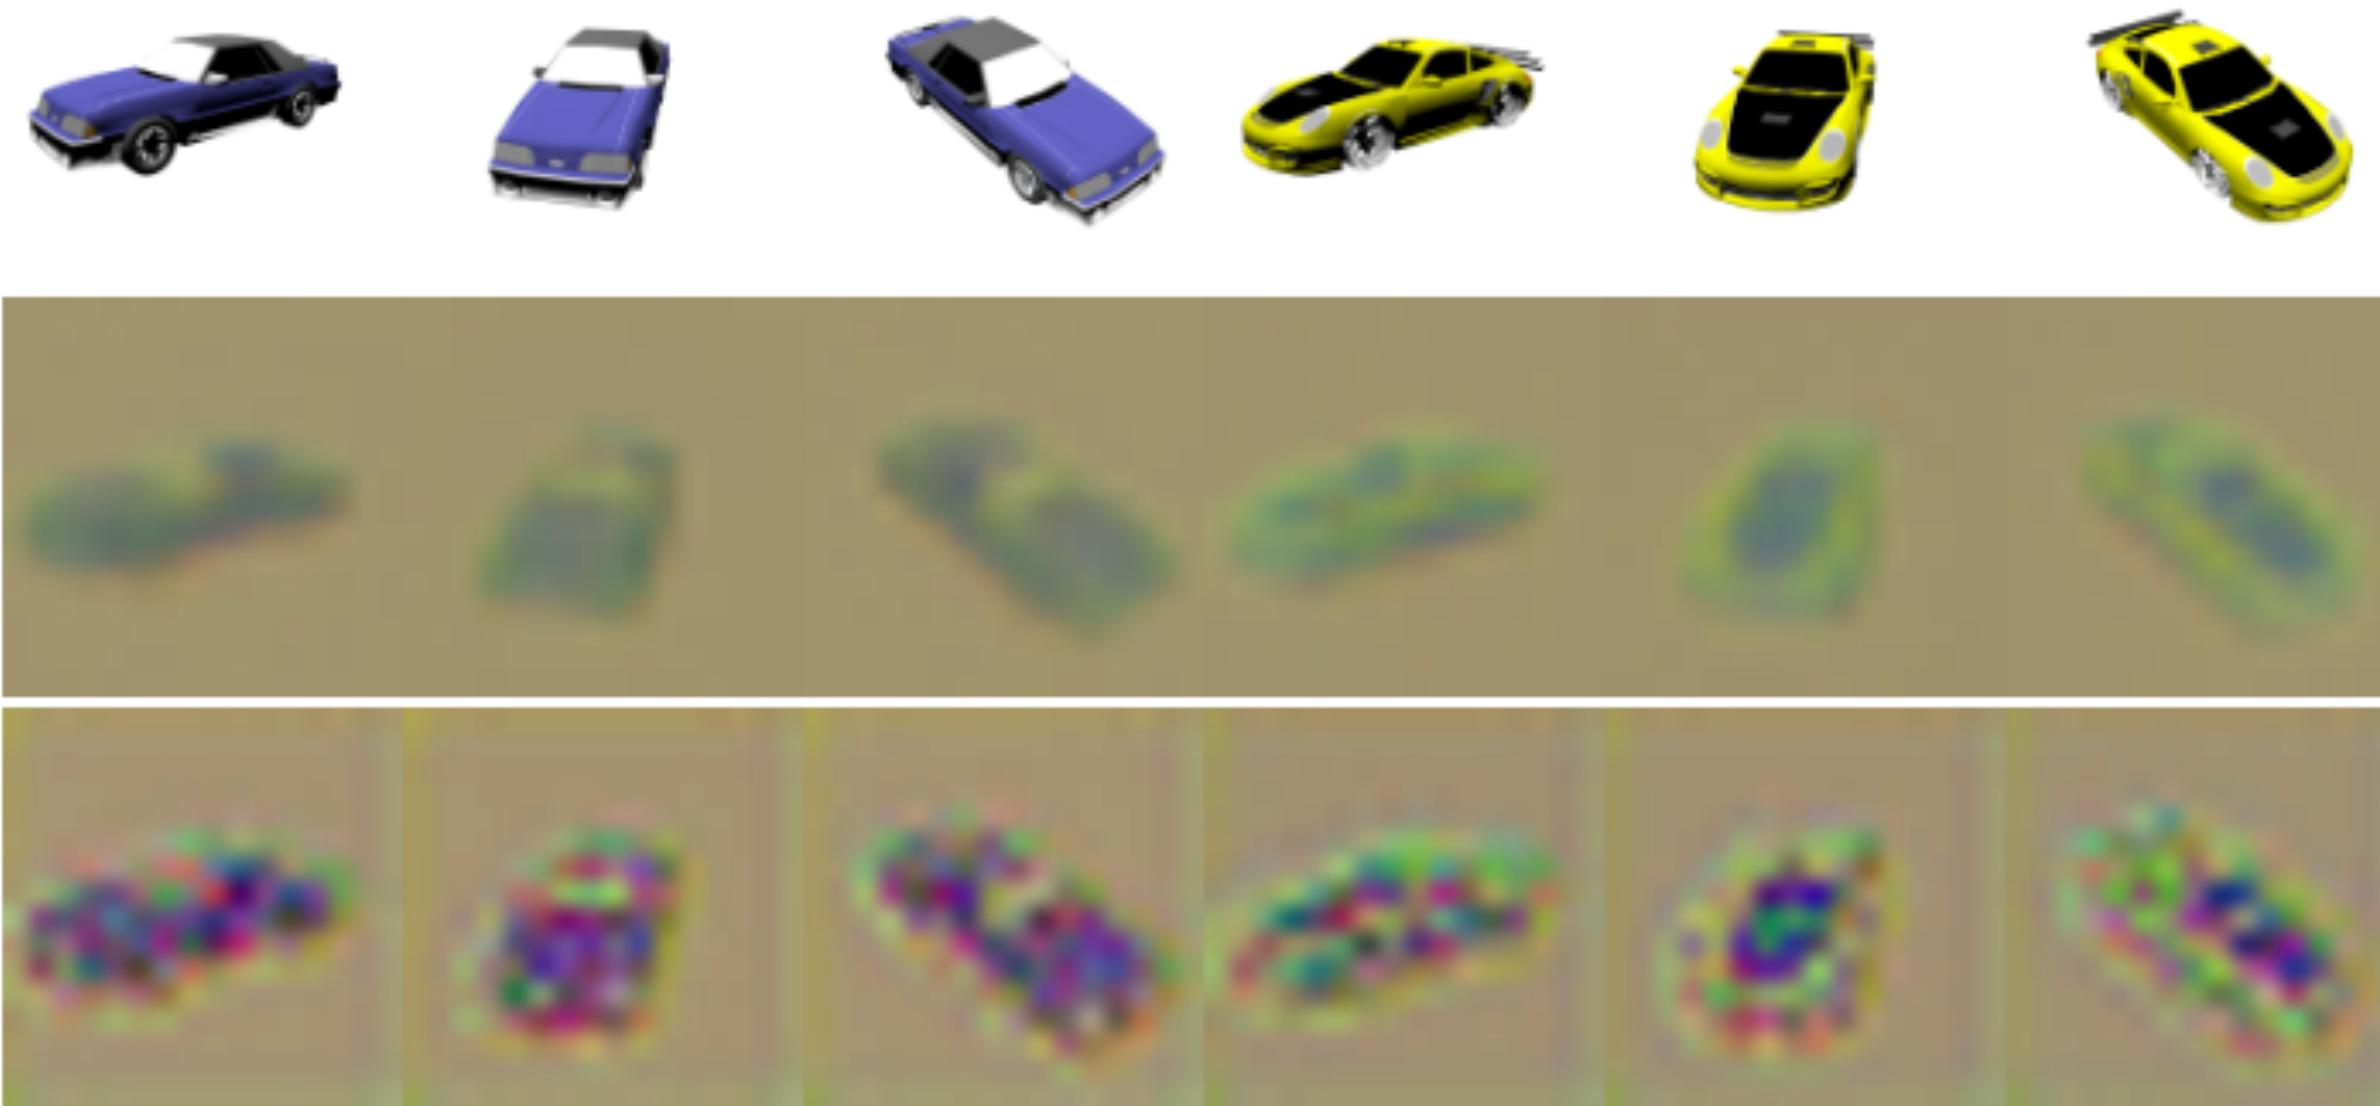


# How to further scale training in a 3D-aware latent space?



# Results

# 3D-aware autoencoder



**Figure 3. Latent space comparison.** Top: ground truth image. Middle: latent image obtained with the 3D-aware encoder. Bottom: latent image obtained with the baseline encoder. Qualitative results show that our 3D-aware encoder better preserves 3D consistency and geometry in the latent space.

# Renderings

Experiment	Latent Space	Micro-Planes	Macro-Planes	Train scenes	Exploit scenes
Ours-Micro	✓	✓	✗	26.52	26.95
Ours-Macro	✓	✗	✓	25.67	26.10
Tri-Planes-Macro (RGB)	✗	✗	✓	27.84	28.00
Tri-Planes (RGB)	✗	✓	✗	<b>28.24</b>	28.40
Ours-No-Prior	✓	✓	✓	27.72	28.13
Ours	✓	✓	✓	28.05	<b>28.48</b>

Table 2. **Quality comparison.** Average PSNR demonstrated by our method with a comparison to Tri-Planes and ablations of our pipeline. All metrics are computed on never-seen test views. Here, we consider  $N_{\text{train}} = 500$ ,  $N_{\text{exploit}} = 100$ , and  $M = 50$ . For compute constraints, Tri-Planes metrics are averaged on 50 scenes.

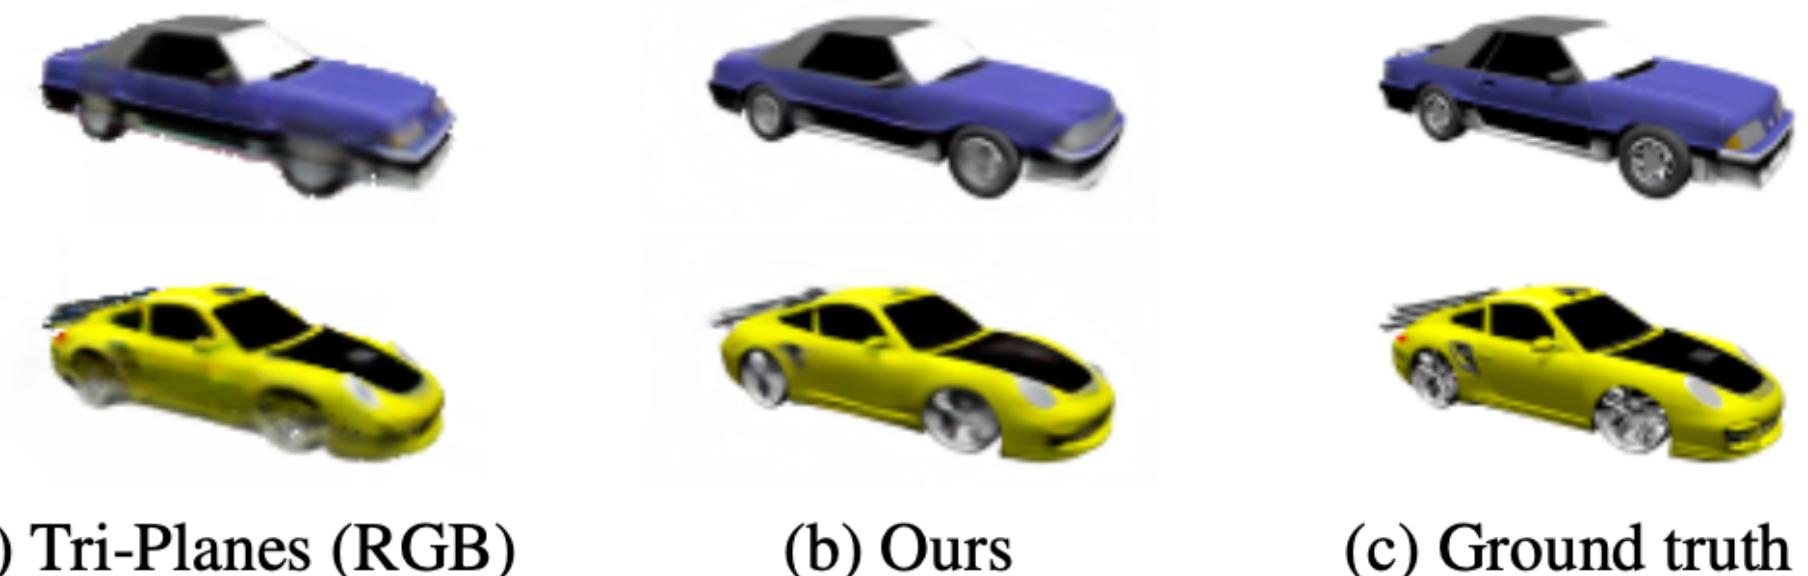


Figure 8. **Visual comparison.** Visual comparison of novel view synthesis quality for our method and Tri-Planes (RGB).

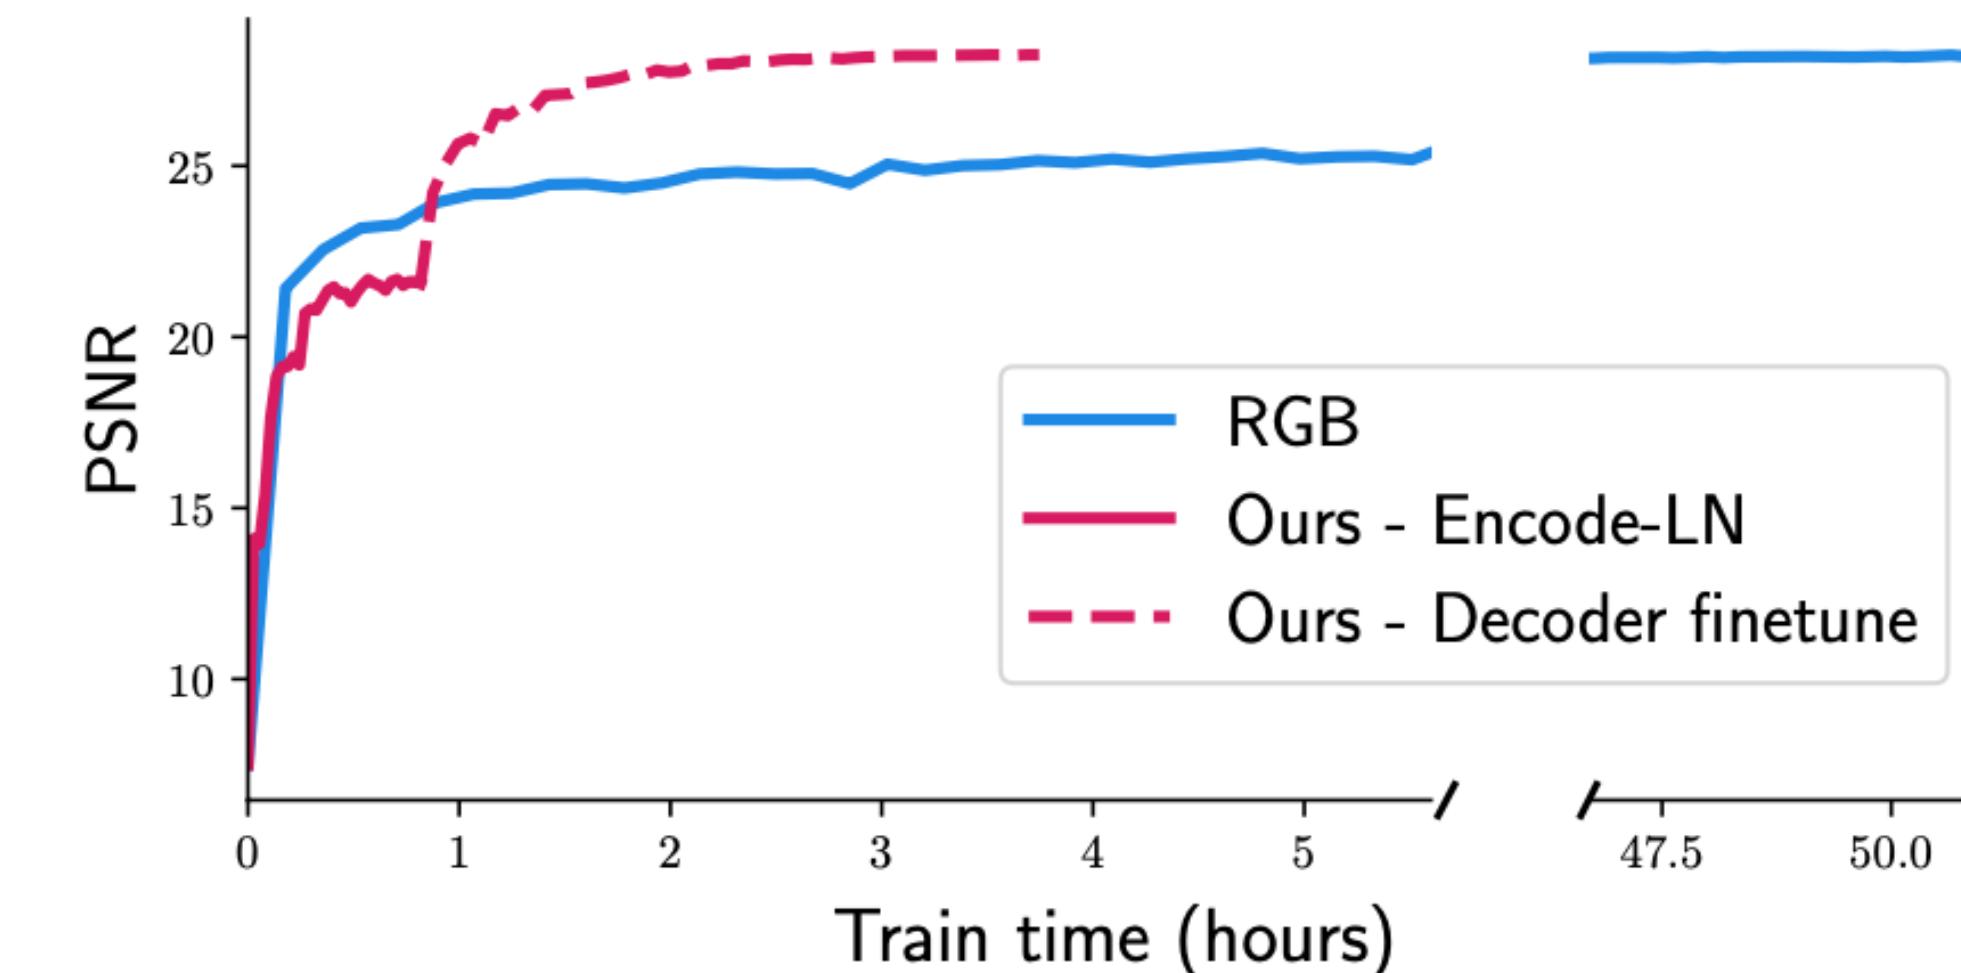


Figure 6. **Quality evolution.** Evolution of the average test-view PSNR demonstrated in the exploit phase of our method compared to RGB Tri-Planes ( $N_{\text{exploit}} = 100$ ). Our method achieves comparable quality in less training time.

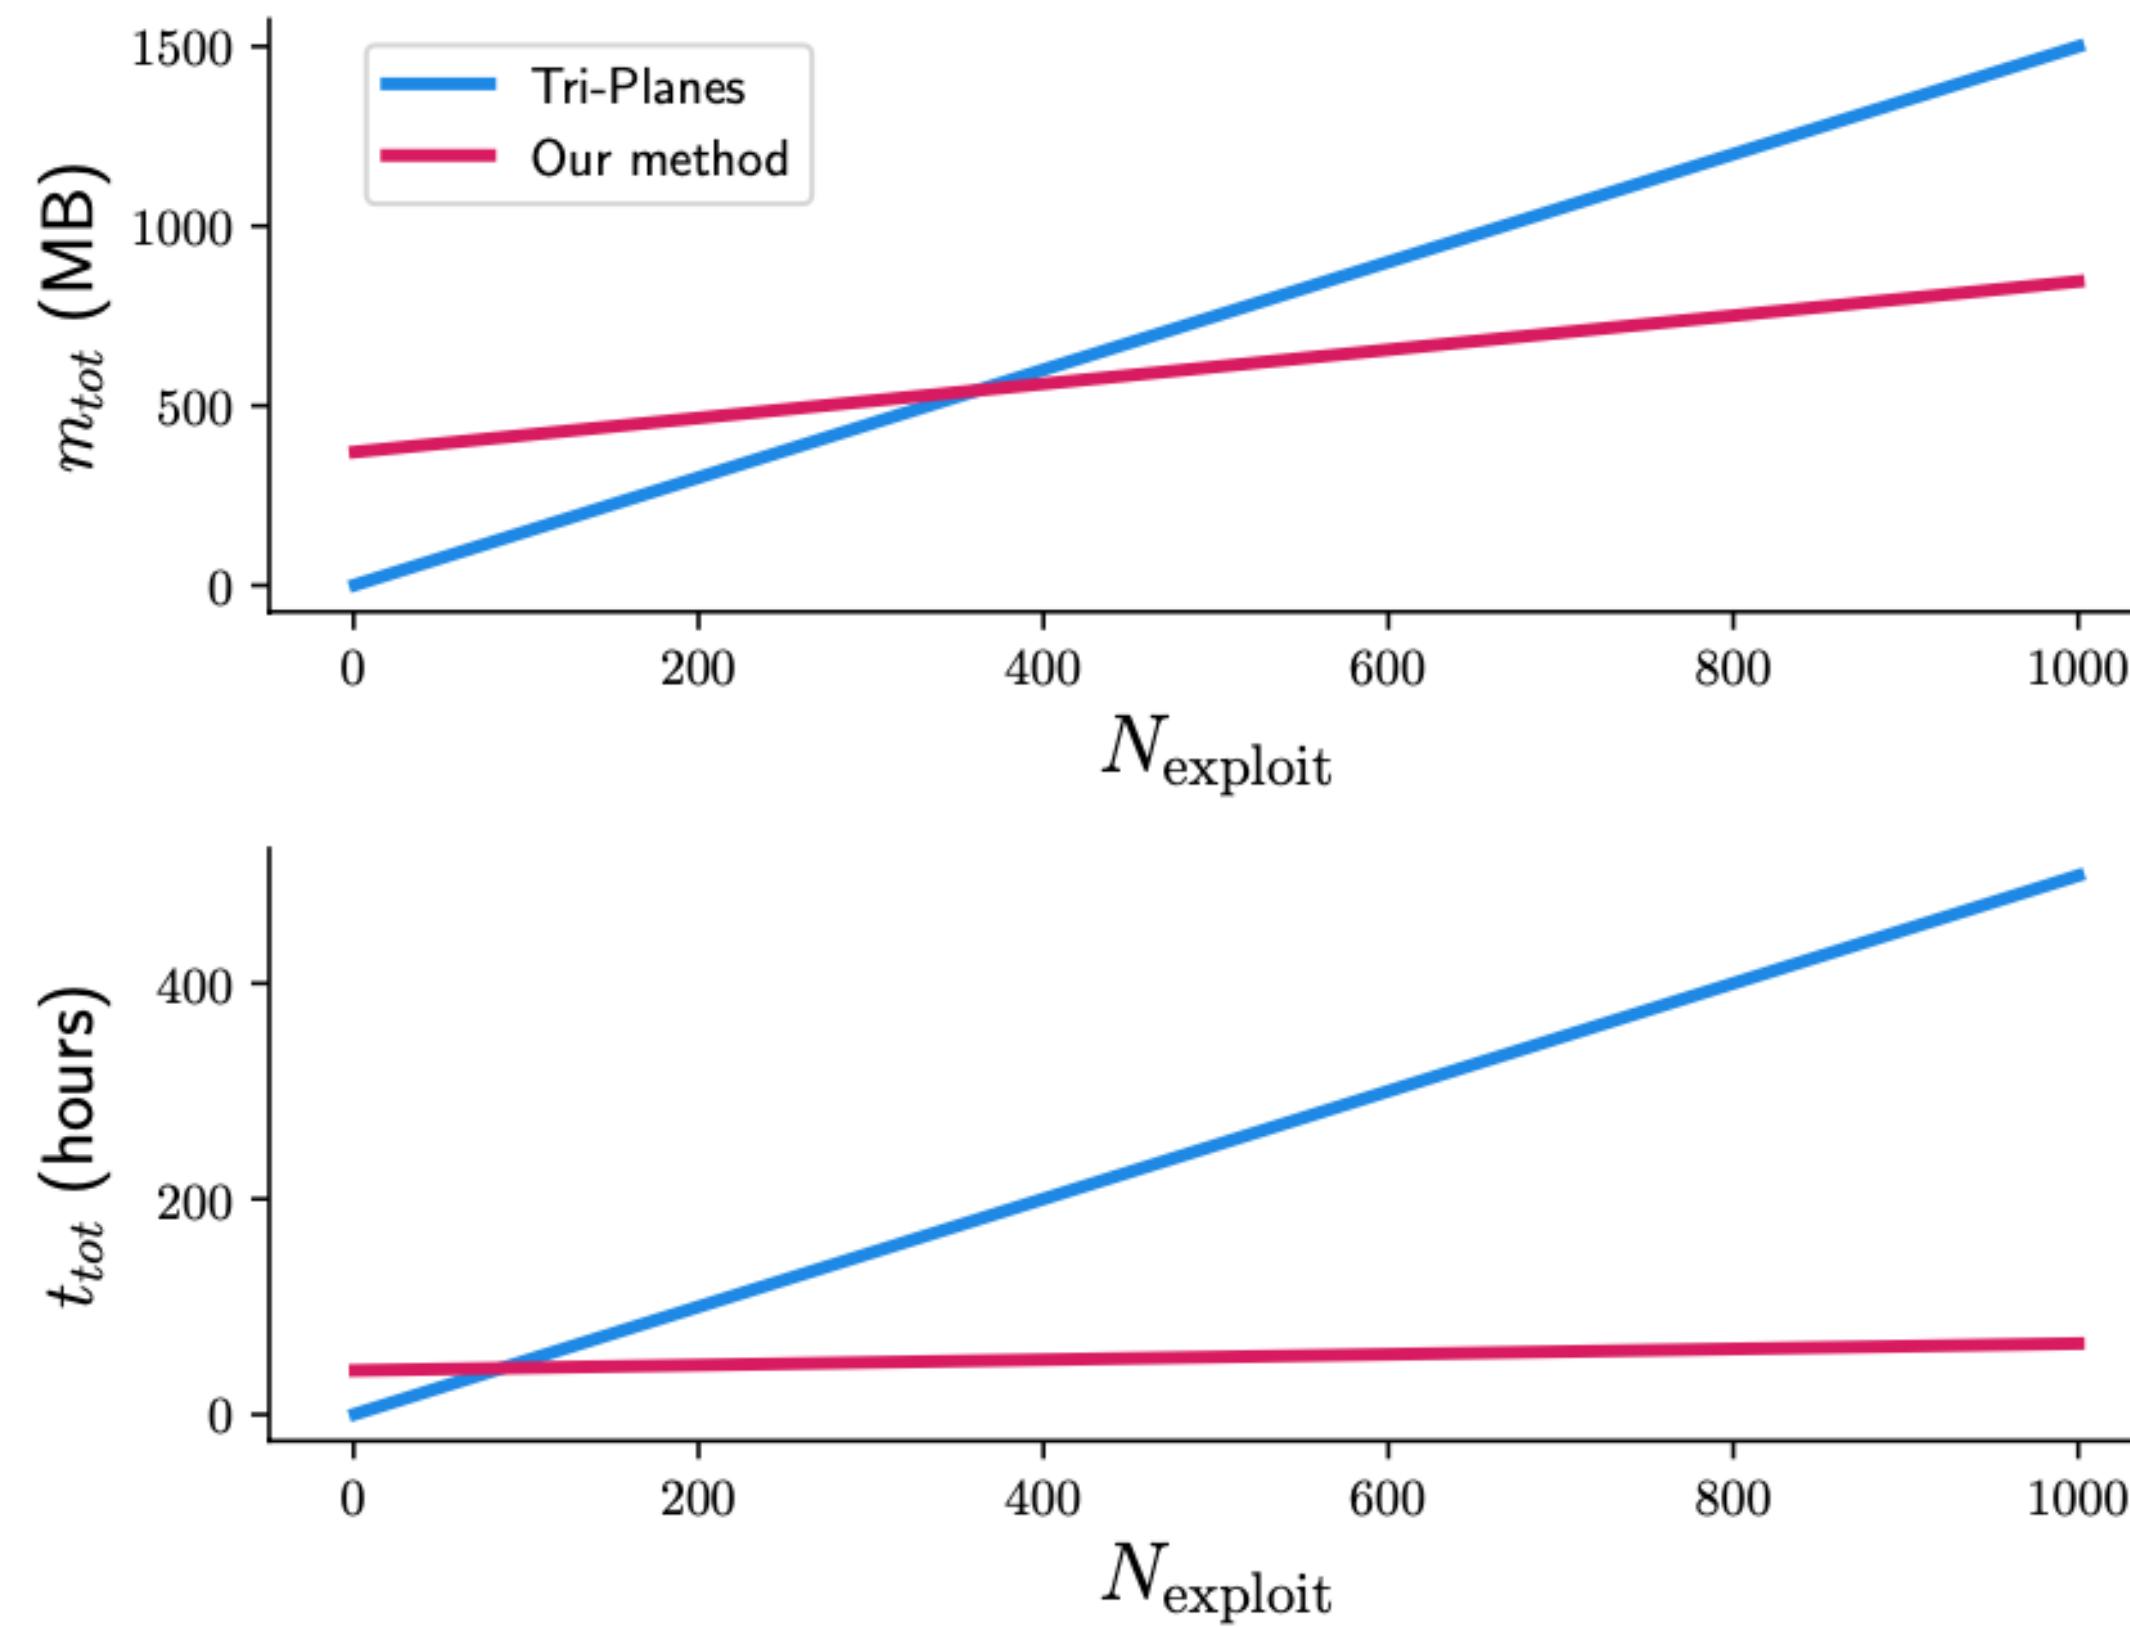
# Resource costs

	$t_{\text{scene}}$ (min)	$t_{\text{eff}}^{\text{scene}}$ (min)	$m_{\text{scene}}$ (MB)	$m_{\text{scene}}^{\text{eff}}$ (MB)	Rendering Time (ms)	Rendering Resolution
Encoder	—	—	0	0.13	—	—
Decoder	—	—	0	0.19	9.7	$128 \times 128$
Tri-Planes (RGB)	32	32	1.5	1.5	23.3	$128 \times 128$
Our method	2	4.5	0.48	0.84	11.0	$128 \times 128$

Table 1. **Cost comparison.** Per scene cost comparison with Tri-Planes trained in the image space. Here, we consider  $N_{\text{train}} = 500$ ,  $N_{\text{exploit}} = 1000$ ,  $t_{\text{EC}} = 40$  hours,  $M = 50$ ,  $F^{\text{mac}} = 22$ . Our method reduces the effective training time by 86% per scene, and the effective memory cost by 44% per scene.

$$t_{\text{scene}}^{\text{eff}} = \frac{t_{\text{EC}}}{N_{\text{exploit}}} + t_{\text{scene}} \quad m_{\text{scene}}^{\text{eff}} = \frac{m_{\text{EC}}}{N_{\text{exploit}}} + m_{\text{scene}}$$

# Resource costs



**Figure 7. Cost evolution.** Total memory and train time evolution when scaling the number of trained scenes  $N_{\text{exploit}}$ . The entry training cost  $t_{EC}$  and memory costs  $m_{EC}$  are taken into account. Our method demonstrates more favorable scalability properties as compared to Tri-Planes (RGB).

# Exploring 3D-aware Latent Spaces for Efficiently Learning Numerous Scenes

Antoine Schnepf<sup>\*1,3</sup>, Karim Kassab<sup>\*1,2</sup>,  
Jean-Yves Franceschi<sup>1</sup>, Laurent Caraffa<sup>2</sup>, Flavian Vasile<sup>1</sup>, Jeremie Mary<sup>1</sup>,  
Andrew Comport<sup>3</sup>, Valérie Gouet-Brunet<sup>2</sup>

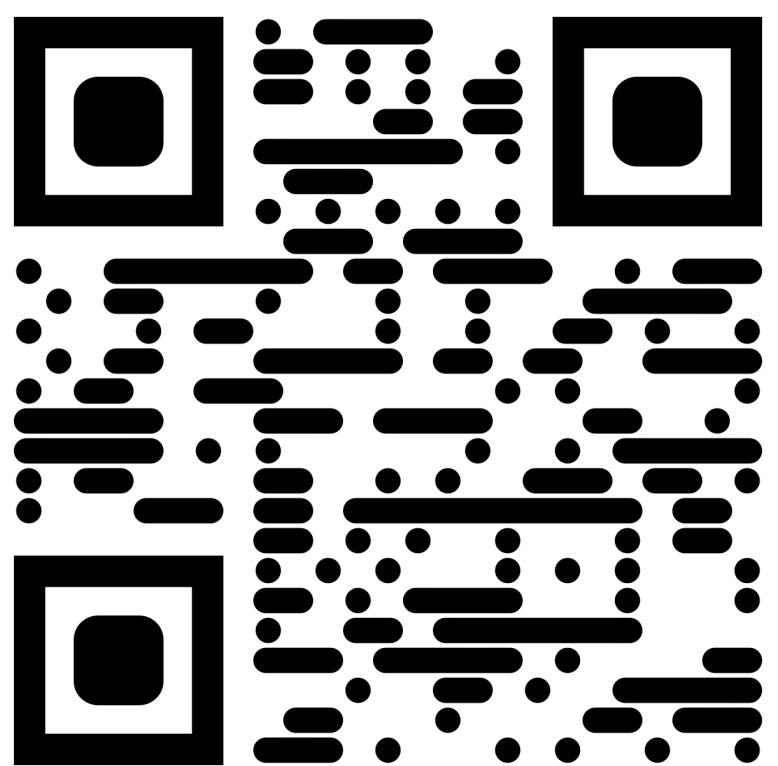
Accepted at 3DMV-CVPR workshop

<sup>\*</sup> Equal Contributions

<sup>1</sup> Criteo AI Lab, Paris, France

<sup>2</sup> LASTIG, Université Gustave Eiffel, IGN-ENSG, F-94160 Saint-Mandé

<sup>3</sup> Université Côte d'Azur, CNRS, I3S, France



3da-ae.github.io